

# Homework #4: Stitching Images into Panoramas

**Assigned: Wednesday, October 16**

**Due: Wednesday, October 30**

The goal of this assignment is to write a simple photo panorama stitcher. You will take four or more photographs and create a panoramic image by computing homographies, projective warping, resampling, and blending the photos into a seamless output image. See the lecture slides for more description of these steps. Specifically, implement the following steps:

## 1. Take Images

Use either some of the images you took for HW #1 or another set of images. One possible set of images is available in [testimages.zip](#). Use at least **four** images to create your panorama. You can assume the input images are named 1.jpg, 2.jpg and so on where the images are taken from left to right in increasing order (i.e., 1.jpg overlaps with 2.jpg and that 1.jpg is to the left of 2.jpg in the original scene). Be sure that consecutive images overlap at least 30%.

## 2. Compute Feature Points

In this step you need to detect feature points and find their correspondences in overlapping pairs of images. To do this automatically use the SIFT keypoint detector. Information on the SIFT detector and descriptor is given in the text in Sections 4.1.1 and 4.1.2. Here is a description of how to use this provided software:

- i. Download the '[hw4SiftCode.zip](#)' zip file to your local machine from the hw4 directory and then unzip it to a local directory. Be sure all files extracted, including `siftWin32.exe`, `sift`, `sift.m`, `match.m`, `isrgb.m`, `appendimage.m` and `showkeys.m`, are put in the *current* working folder. If you only put the unzipped `siftDemoV4` folder in the current working folder, MATLAB will not be able to find the code.
- ii. Read the 'README' file to understand how to use the functions, including `sift` and `match`. For example, the following two lines of code show you how to use and display the results of the function `sift`:
 

```
[image, descripts, locs] = sift('foo.jpg');
showkeys(image, locs);
```

`showkeys` will display a figure showing the feature points detected in image `foo.jpg`. In the figure the tail of the arrow is the position of the feature point, the direction of the arrow indicates the dominant orientation, and the length indicates its scale.
- iii. The `match.m` file lets you access the *index* of matched feature points as follows:
 

```
function [num, match, loc1, loc2] = match(image1, image2)
```

 where the first output parameter is the number of matched feature points between `image1` and `image2`, the second output parameter is the index of matched feature points in `image1` and `image2`, and the third and fourth parameters are the feature point locations (row, column, scale, orientation) in `image1` and `image2` respectively.
- iv. Call the `match` function using:
 

```
[num, matchIndex, loc1, loc2] = match(image1, image2);
```

 where, for example, `image1` is `foo.jpg` and `image2` is `bar.jpg`. If you want to change the default parameter that determines when two points match, modify the

- variable called `distRatio` in `match.m`
- v. Extract the coordinates of the matched feature points in `image1` and `image2` using
 

```
im1_ftr_pts = loc1(find(matchIndex > 0), 1 : 2);
im2_ftr_pts = loc2(matchIndex(find(matchIndex > 0)), 1 : 2);
```

 In vector `matchIndex` elements that do *not* equal 0 are feature points in `image1` that have matches in `image2`, and the values represent the positions of the corresponding feature points in `image2`. The first column of `im1_ftr_pts` keeps the row index in `image1` of every matched feature point, and the second column of `im1_ftr_pts` keeps the column index of every matched feature point. Similarly for `im2_ftr_pts`. For example, if `match` returns the `matchIndex` value `[0 0 8 0 0 15 0]` it means there are two pairs of matched features. The first pair is the 3<sup>rd</sup> feature point in `image1` and the 8<sup>th</sup> feature point in `image2`. The second matched pair is the 6<sup>th</sup> feature point in `image1` and the 15<sup>th</sup> feature point in `image2`.
  - vi. The previous step computes two  $n \times 2$  matrices, `im1_ftr_pts` and `im2_ftr_pts`, which hold the (*row*, *column*) locations of all the matched feature points in `image1` and `image2`, respectively, where  $n$  is the number of matched points.

### 3. Compute Homographies

Using the point correspondences found in the previous step, you next need to recover the parameters of the transformation between each pair of overlapping images. In our case, the transformation is a homography:  $p' = Hp$ , where  $H$  is a  $3 \times 3$  matrix with 8 degrees of freedom (lower-right corner is a scaling factor and can be set to 1). One way to compute the homography is via a set of  $(p', p)$  pairs of corresponding feature points taken from the two images, as computed in the previous step. Given these correspondences, write a function of the form:

```
H = calcH(im1_ftr_pts, im2_ftr_pts)
```

where `im1_pts` and `im2_pts` are  $n \times 2$  matrices holding the  $(x, y)$  locations of  $n$  point correspondences from the two images, and  $H$  is the recovered  $3 \times 3$  homography matrix. In order to compute the entries in the matrix  $H$ , you need to set up a linear system of  $n$  equations (i.e., a matrix equation of the form  $Ah = b$  where  $h$  is a column vector holding the 8 unknown entries of  $H$ ). If  $n = 4$ , the system can be solved using a standard technique. However, with only four points, homography recovery is very unstable and prone to noise. Therefore more than 4 correspondences should be used to produce an overdetermined system, which can be solved using least-squares. In MATLAB, this can be performed using the MATLAB “\” operator (see `mldivide` for details). Information on homographies is given in the textbook in Sections 2.1.2, 6.1.2, 6.1.3, and 9.1.1.

Here is some sample code to compute the homography matrix  $H$ :

```
function H = calcH(p1, p2)
% Assume we have two images called '1' and '2'
% p1 is an n x 2 matrix containing n feature points, where each row
% holds the coordinates of a feature point in image '1'
% p2 is an n x 2 matrix where each row holds the coordinates of a
% corresponding point in image '2'
% H is the homography matrix, such that
% p1_homogeneous = H * [p2 ones(size(p2, 1), 1)]'
% p1_homogeneous contains the transformed homogeneous coordinates of
% p2 from image '2' to image '1'.
```

```

n = size(p1, 1);
if n < 4
    error('Not enough points');
end
A = zeros(n*3,9);
b = zeros(n*3,1);
for i=1:n
    A(3*(i-1)+1,1:3) = [p2(i,:),1];
    A(3*(i-1)+2,4:6) = [p2(i,:),1];
    A(3*(i-1)+3,7:9) = [p2(i,:),1];
    b(3*(i-1)+1:3*(i-1)+3) = [p1(i,:),1];
end
x = (A\b)';
H = [x(1:3); x(4:6); x(7:9)];

```

Unfortunately, not all the corresponding points found by SIFT will be correct matches. To eliminate outliers, you'll need to implement the **RANSAC algorithm** as part of the procedure for computing an homography. That is, select four pairs of points randomly from those returned by SIFT, compute **H** from these four pairs of points, and then see how many of the *other* pairs of points agree (i.e., a point projects very near its corresponding point in the pair. You may set a distance threshold to decide whether two points are close enough.) Repeat this many times (say 100) and choose the homography **H**<sup>\*</sup> with the most agreement. Finally, compute the final homography using *all* the points that agree with **H**<sup>\*</sup>. You should implement RANSAC only at the end *after* you get a complete version working that does *not* use RANSAC. (For more information on RANSAC, see Section 6.1.4 in the textbook.)

One more thing: You should compute the homography matrices so as to warp each image *directly* to the reference image plane. For example, assume you have images 1, 2, 3, 4 and 5, and 3 is the reference image. You have homography matrices warping 1→2, 2→3, 4→3, and 5→4. You should create a SINGLE homography matrix for each non-reference image by composing one or more homography matrices so that you can warp directly to the reference plane rather than warping to one plane and then warping the result to another. Here is how to do this: Assume you've got the homography matrix to warp 5→4 called **H\_54** which is a  $3 \times 3$  matrix used as follows:  $p_4 = \mathbf{H}_{54} * p_5$ , and 4→3: **H\_43** which is a  $3 \times 3$  matrix used as:  $p_3 = \mathbf{H}_{43} * p_4$ , where  $p_5$ ,  $p_4$  and  $p_3$  are the *homogeneous coordinates* of corresponding points in images 5, 4 and 3, respectively, and each is a  $3 \times 1$  column vector. Below are two possible ways to map point  $p_5$  in image 5 into its coordinates in image 3:

First method:

- (a) Warp point  $p_5$  to image 4 using **H\_54** by computing  $p_4 = \mathbf{H}_{54} * p_5$ , where  $p_4$  is the corresponding point in image 4.
- (b) Warp point  $p_4$  to image 3 using **H\_43** by computing  $p_3 = \mathbf{H}_{43} * p_4 = \mathbf{H}_{43} * \mathbf{H}_{54} * p_5$
- (c) Finally, convert  $p_3$  from homogeneous coordinates to Cartesian coordinates using  $p_3 = p_3 ./ p_3(1,3)$ . This results in the *Cartesian coordinates* where point  $p_5$  should be in image 3.

Second method:

Alternatively, first compute  $\mathbf{H}_{53} = \mathbf{H}_{43} * \mathbf{H}_{54}$ . Then compute  $\mathbf{p}_3 = \mathbf{H}_{53} * \mathbf{p}_5$ , and convert from homogeneous coordinates to Cartesian coordinates using  $\mathbf{p}_3 = \mathbf{p}_3 ./ \mathbf{p}_3(1, 3)$  to get the *Cartesian coordinates* where point  $\mathbf{p}_5$  should be in image 3.

The above two methods will give the similar results though the second method is much better because it is easier to implement, introduces fewer resampling errors, and is faster. So, you should implement the second method.

#### 4. Warp Images

The next step is to warp all the input images into the output image. For large field of view panoramas, it is common to use a cylindrical surface for this purpose. When only a few images are used covering a relatively small field of view, however, we can more simply warp all the images onto a plane defined by one of the input images, which we'll call the **reference image**. We call the result image a **flat panorama**.

- i. Before warping each of the images, first compute the size of the output flat panorama image by computing the range of warped image coordinates for each input image. Do this by mapping the coordinates of the four corners (i.e., the top-left, top-right, etc.) from *each* source image using *forward warping* to determine its coordinates in the output image. Then compute the minimum\_r, minimum\_c, maximum\_r, and maximum\_c coordinates to determine the size of the output image. That is, the width of output image should be  $\text{maximum\_c} - \text{minimum\_c} + 1$ . Finally, compute r\_offset and c\_offset values that specify the offset of the origin of the reference image relative to the origin of the output image. That is,  $\text{r\_offset} = 1 - \text{minimum\_r}$ , if we define the top-left corner of the image as (1,1). You may want to use the MATLAB `meshgrid` function to generate a 2D array of points for the output panorama image. Initialize all pixels to black (0).
- ii. Step i uses forward warping to determine the size for the output image. However, if you use forward warping to map every pixel from each source image, there will be holes (i.e., some pixels in the output image will not be assigned an RGB value from *any* source image, and remain black) in the final output image. Therefore, we need to use *inverse warping* to map each pixel in the output image (that you created in the previous step) into the planes defined by the corresponding source images. Use **bilinear interpolation** to compute the color values in the warped image. The MATLAB function `interp2` (with argument 'linear' to do bilinear interpolation) can be used for this purpose, though it may well be simpler to just implement this yourself. Here's sample code that inverse warps a pixel (at coordinates (r, c)) in the output image onto the plane defined by the 1<sup>st</sup> image and then sets the color value appropriately:

```
% Inverse warping the coordinates. Here it is assumed that 3 is the
% reference image and point (r, c) is mapped to the 1st image.
```

```
M = H_31*[r c 1]';
```

```
% Let (x,y) be the corresponding warped coordinates of (r,c) in the
% coordinate system of the 1st image
```

```
x = M(1);
```

```
y = M(2);
```

```
% Set the color value of the output image at (r,c)
outputImage( r, c, : ) = image1( x, y, : );
```

The above procedure has to be repeated for all the pixels in the output image. Note that part of the problem is to figure out to which of the source images the pixel  $(r, c)$  in the output image gets mapped to. In the above example it is assumed that the point  $(r, c)$  maps to (in other words, lies in) the 1<sup>st</sup> image. It is also assumed that the warped coordinates  $(x, y)$  are integers, which may not always be the case (in fact it rarely is the case). If the values for  $x$  and  $y$  are decimals, then use the `interp2` command, as mentioned earlier, to interpolate the color value at  $(x, y)$  from its 4 surrounding pixels. For example, to get the red component value at  $(x, y)$  do:

```
R = interp2(1:size(image1,2),1:size(image1,1),image1(:,:,1), y, x);
```

Be sure to first convert `image1` to type `double` before calling `interp2` because `uint8` images may cause an error.

Also be aware that several images may sometimes overlap at a given pixel in the output image, and we will deal with this in the next step (see below).

NOTE: You can do inverse warping pixel by pixel in the output image, but it will take a considerable amount of time if dealing with large images. In order to make your program run efficiently, it is best if you can vectorize your code using `interp2` and `meshgrid`. The fourth and fifth arguments for `interp2` can take a vector or matrix as input so you will only need to call `interp2` three times (once for each RGB color channel) to interpolate all the pixels. Type ‘`help interp2`’ and ‘`help meshgrid`’ to view the help documents for more details.

## 5. Blend the Images

Given the warped images and their relative displacements, the final step is to blend overlapping pixel color values in such a way as to hide seams. One simple way to do this, called **feathering**, is to use weighted averaging of the color values to blend overlapping pixels (see Section 9.3.2 in the text). To do this, use an ‘alpha channel’ where the value of alpha for an image is 1 at its center pixel and decreases linearly to 0 at all the border pixels. You can use the Matlab function `bwdist` with argument ‘`euclidean`’ and applied to a binary image that has 1s in the first and last rows, and first and last columns, to compute this distance map and then convert it into weights in the range 0 – 1. Use these alpha values as follows to compute the color at a pixel where at least two images overlap in the output image. For example, suppose there are 2 images,  $I_1, I_2$ , overlapping in the output image; each pixel  $(x, y)$  in image  $I_i$  is represented as  $I_i(x, y) = (\alpha_i R, \alpha_i G, \alpha_i B, \alpha_i)$  where  $(R, G, B)$  are the color values at the pixel. Compute the pixel value of  $(x, y)$  in the output image as  $[(\alpha_1 R, \alpha_1 G, \alpha_1 B, \alpha_1) + (\alpha_2 R, \alpha_2 G, \alpha_2 B, \alpha_2)] / (\alpha_1 + \alpha_2)$ . Save your output image as a jpeg file.

Here is some sample code to do simple blending:

```
% input image name
appleImgName = 'apple.jpg';
orangeImgName = 'orange.jpg';
% read images
apple = imread(appleImgName);
```

```

orange = imread(orangeImgName);
% convert to type double
apple = double(apple);
orange = double(orange);
% check to see if the dimensions are the same
if(~((size(apple, 1) == size(orange, 1)) && (size(apple, 2) ==
size(orange, 2))))
    error('Size does not match!');
end
% get dimension information
height = size(apple, 1);
width = size(apple, 2);
channel = size(apple, 3);
% weight vector
weightApple = 1 : -1/(width-1) : 0;
% weight matrix: repeat of weight vector
weightAppleMatrix = repmat(weightApple, size(apple, 1), 1);
% weight vector
weightOrange = 1 - weightApple;
% weight matrix: repeat of weight vector
weightOrangeMatrix = repmat(weightOrange, size(orange, 1), 1);
% create the blending image
blendingImg = zeros(size(apple));
% compute pixel value for blending image
for i = 1 : channel
    blendingImg(:, :, i) = apple(:, :, i) .* weightAppleMatrix +
orange(:, :, i) .* weightOrangeMatrix;
end
% convert output image into uint8 type
blendingImg = uint8(blendingImg);
% display the output image
figure;
imshow(blendingImg);
title('Image blending example: appange? oranple?');
% write the output image to disk
blendingImgName = 'blending.jpg';
imwrite(blendingImg, blendingImgName, 'jpg');

```

## Problems you might face while working with SIFT

- i. Error message: `./sift: Permission Denied`  
 You might get this error if you don't have permission to execute the 'sift' or 'siftWin32' programs. To resolve this issue, change the permissions of these files so that you (as the user) have executable permissions with them.
- ii. Error message: "Error using ==> ctranspose, Transpose on ND array is not defined."  
 Uncomment the lines 26 - 28 in `sift.m`
- iii. SIFT doesn't work on Mac OS.  
 The provided executables run on Linux or Windows machines only.

- iv. Error message: `"/bin/bash: ./sift: No such file or directory ???  
Error using ==> sift at 57 Invalid keypoint file beginning."`  
Include the executable 'sift' file **directly** under the current working folder or MATLAB path,  
or modify the files to enable both read and write permissions.

If you still face problems with SIFT you can try using the SIFT keypoint detector by VLFeat.

## Program Instructions

Your main function should be called 'panorama' and should be written in a file named `panorama.m`. It should be called as follows:

```
output = panorama('path/to/images');
```

where `output` is a 3D matrix of type `uint8` that contains the color panoramic output image.  
Here's how an example run of your function should work:

```
output = panorama('images/');
```

Also create a `.jpg` file of your output image. Note that you can write any other supporting functions and files as needed. Skeleton code that includes *some* of the required steps is given in the provided file called `panorama-skeleton.m`.

## Extra Credit

Instead of mapping the images onto a plane, map them onto a cylinder by using cylindrical projection as described in the lecture notes. See Section 9.1.6 in the text for more information. To interactively view your resulting panorama you can use one of the existing viewers listed on the homework page.

## Hand-In Instructions:

Put all the code you wrote (as one or more `.m` files), the input images you used (if they are not any of the ones provided), the output panorama image(s) you created (as `.jpg` files), and a `README.txt` file that contains comments on how to execute your program, any relevant implementation notes including parameter values used, and any extra work beyond the basic requirements, if you did any, into a single zip file called `lastname.firstname.hw4.zip` and submit this *one* file to Moodle.