

# Today

---

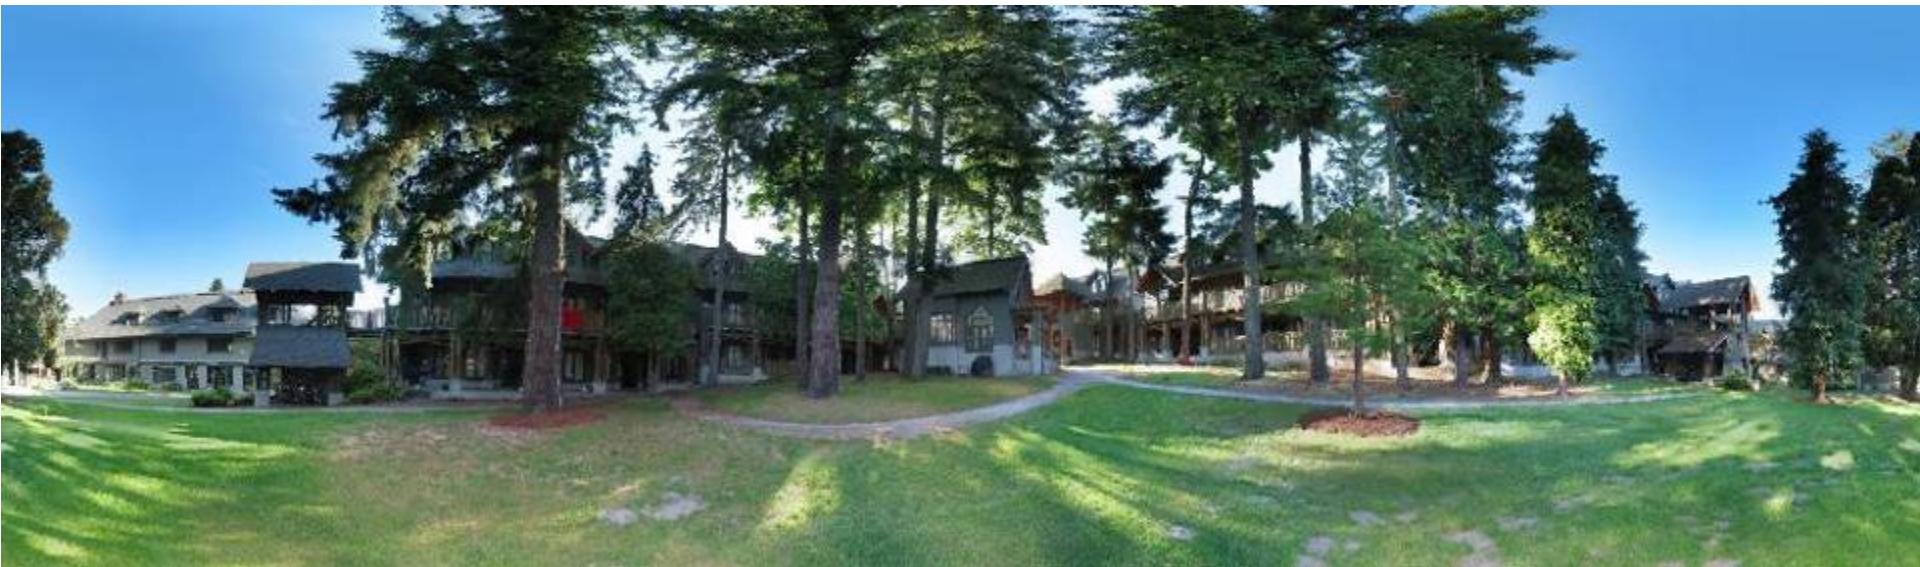
## Homographies, Mosaics, Panorama (some slides from Snavely)

- Today's Readings

Szeliski and Shum paper

<http://www.acm.org/pubs/citations/proceedings/graph/258734/p251-szeliski/>

Hartley Zisserman (Ch 3—4)



Full screen panoramas (cubic): <http://www.panoramas.dk/>

Mars: [http://www.panoramas.dk/fullscreen3/f2\\_mars97.html](http://www.panoramas.dk/fullscreen3/f2_mars97.html)

2003 New Years Eve: <http://www.panoramas.dk/fullscreen3/f1.html>

# Why Mosaic?

---

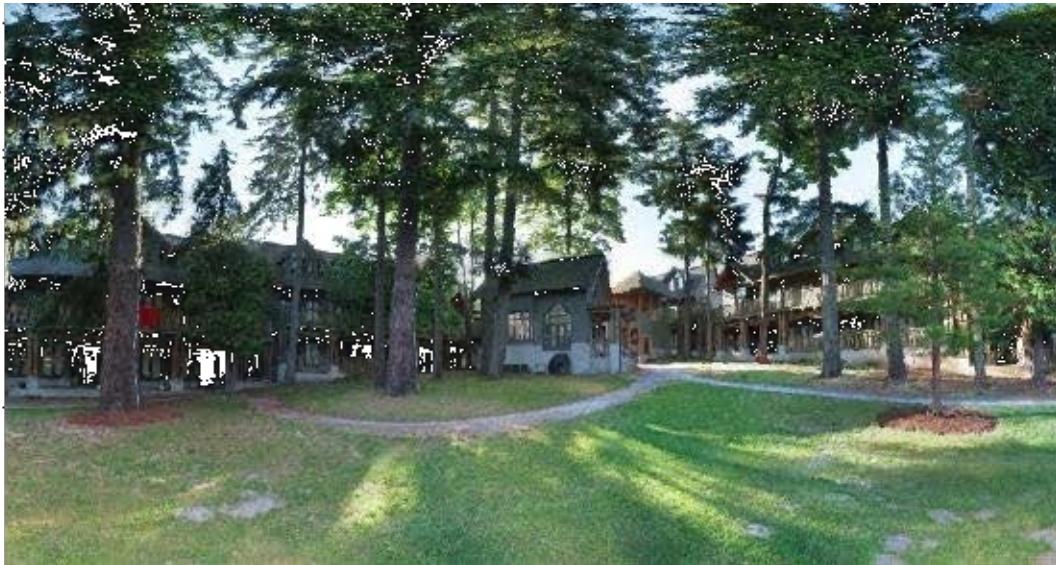
- Are you getting the whole picture?
  - Compact Camera FOV =  $50 \times 35^\circ$



# Why Mosaic?

---

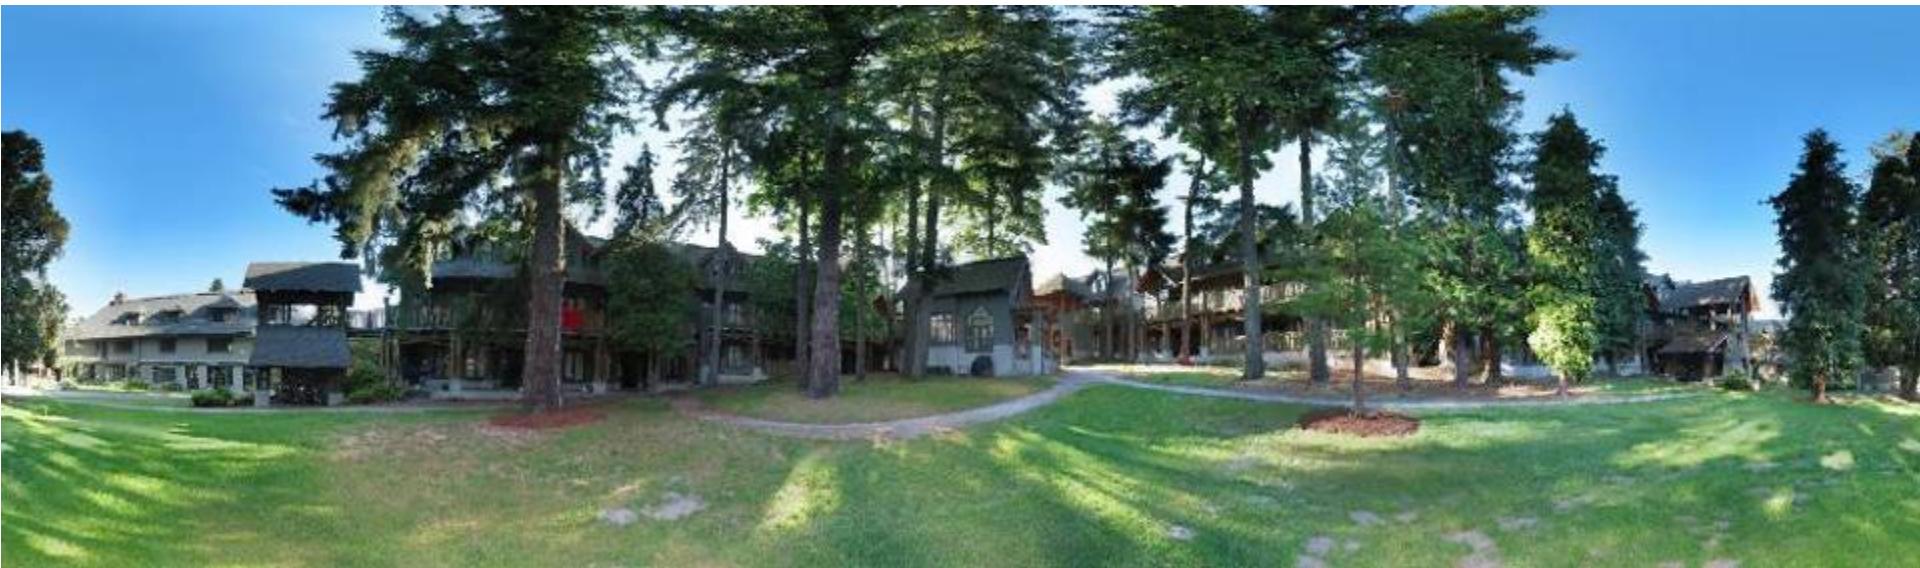
- Are you getting the whole picture?
  - Compact Camera FOV =  $50 \times 35^\circ$
  - Human FOV =  $200 \times 135^\circ$



# Why Mosaic?

---

- Are you getting the whole picture?
  - Compact Camera FOV =  $50 \times 35^\circ$
  - Human FOV =  $200 \times 135^\circ$
  - Panoramic Mosaic =  $360 \times 180^\circ$



# Mosaics: stitching images together

---



Creating virtual wide-angle camera

# Madison Panoramas

---



Hans Werner, 2011

# Madison Panoramas

---



Jean Forde, 2011

# Madison Panoramas

---



George Wanant, 2010

# How to do it?

---

- Basic Procedure
  - Take a sequence of images from the same position
    - Rotate the camera about its optical center

# How to do it?

---

- Basic Procedure
  - Take a sequence of images from the same position
    - Rotate the camera about its optical center
  - Compute transformation between second image and first (somehow)
  - Transform the second image to overlap with first

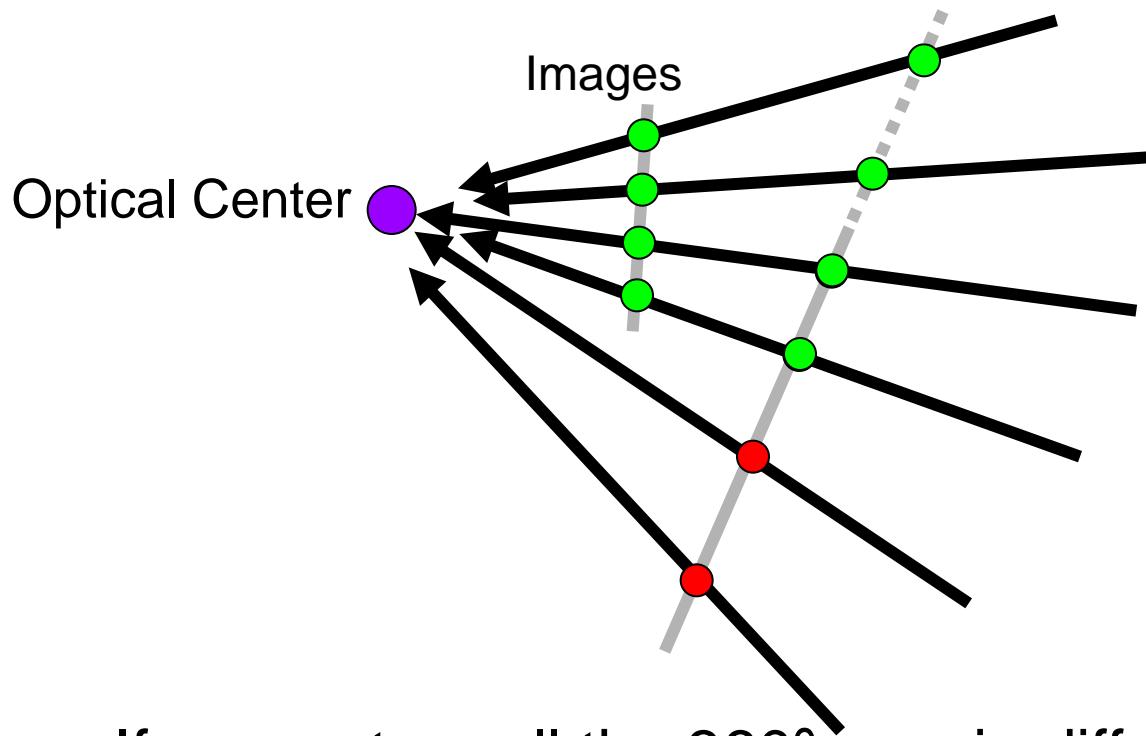
# How to do it?

---

- Basic Procedure
  - Take a sequence of images from the same position
    - Rotate the camera about its optical center
  - Compute transformation between second image and first
  - Transform the second image to overlap with first
  - Blend the two together to create a mosaic
  - If there are more images, repeat

# Geometric Interpretation of Mosaics

---



- If we capture all the  $360^{\circ}$  rays in different images, we can assemble them into a panorama.
- The basic operation is projecting an image from **one plane to another**

# What is the geometric relationship between these two images?

---



# What is the transformation?

---



left on top



right on top



# What is the transformation?

---



left on top



right on top



# What is the transformation?

---



left on top



right on top



Translations are not enough to align the images



# Where to go from affine transformations?

$$\begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix}$$

affine transformation

# Where to go from affine transformations?

---

$$\begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix}$$

affine transformation



This is what happens when we change this row...

# Projective Transformations aka Homographies aka Planar Perspective Maps

$$\mathbf{H} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & 1 \end{bmatrix}$$

Called a *homography*  
(or *planar perspective map*)



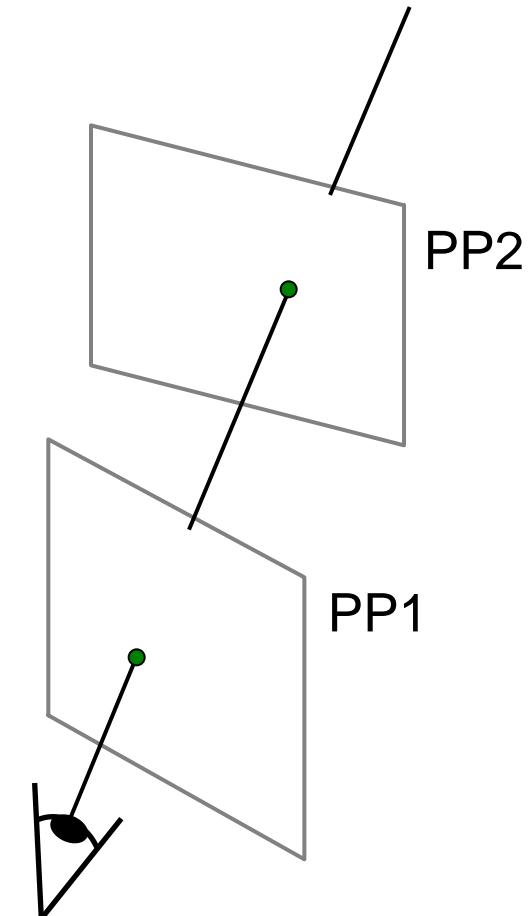
projection of 3D plane

# Homography

- Mapping between any two PPs with the same center of projection
  - rectangle should map to arbitrary quadrilateral

$$\begin{bmatrix} wx' \\ wy' \\ w \end{bmatrix} = \begin{bmatrix} * & * & * \\ * & * & * \\ * & * & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

$\mathbf{p}'$        $\mathbf{H}$        $\mathbf{p}$



To apply a homography  $\mathbf{H}$

- Compute  $\mathbf{p}' = \mathbf{H}\mathbf{p}$  (regular matrix multiply)
- Convert  $\mathbf{p}'$  from homogeneous to image coordinates

# Image warping with homographies

---

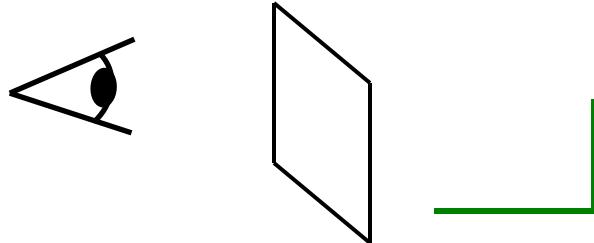


image plane in front

# Image warping with homographies

---



$$H_1 \rightarrow$$

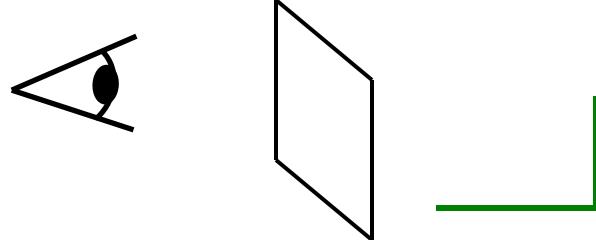
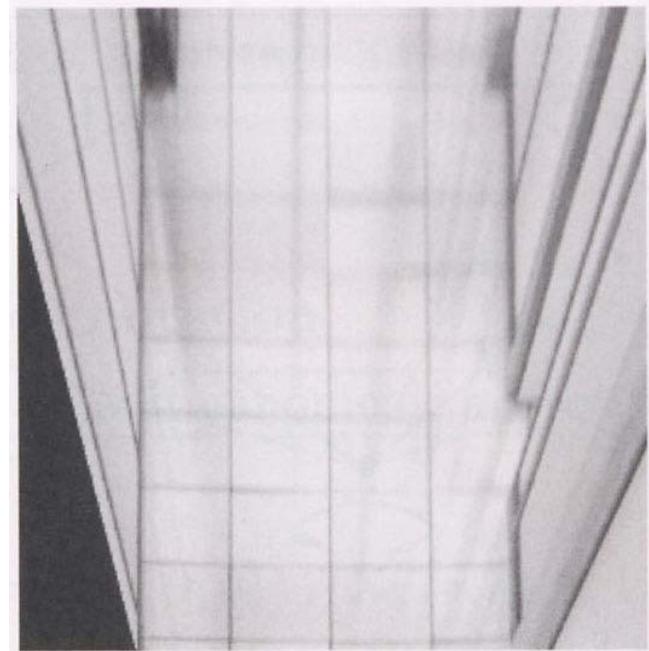


image plane in front

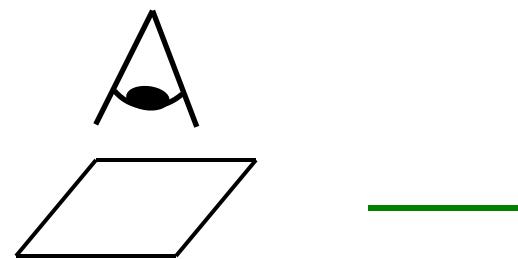
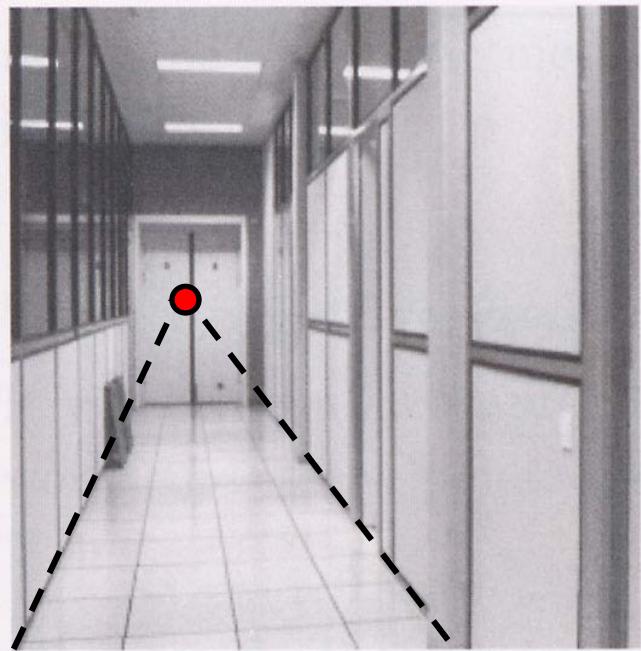


image plane below

# Image warping with homographies



$$H_1 \rightarrow$$

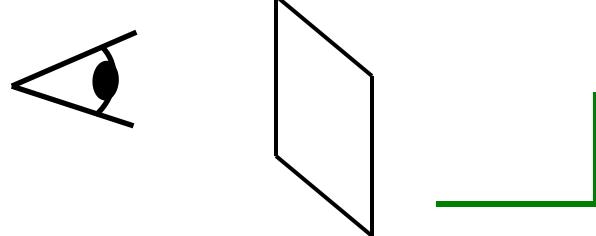
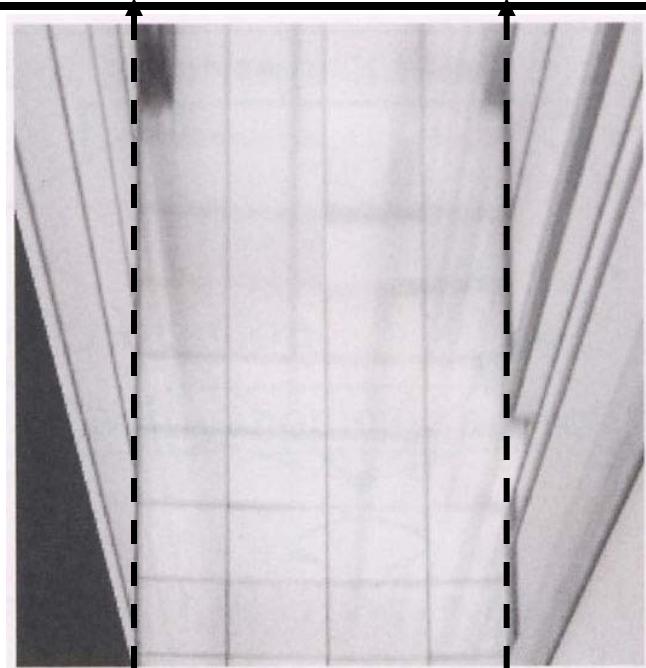


image plane in front

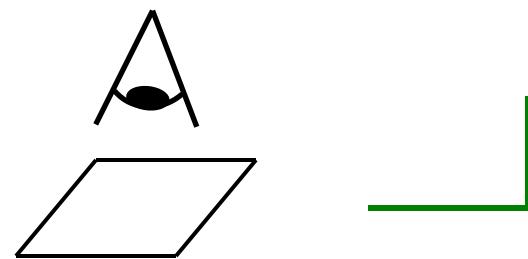
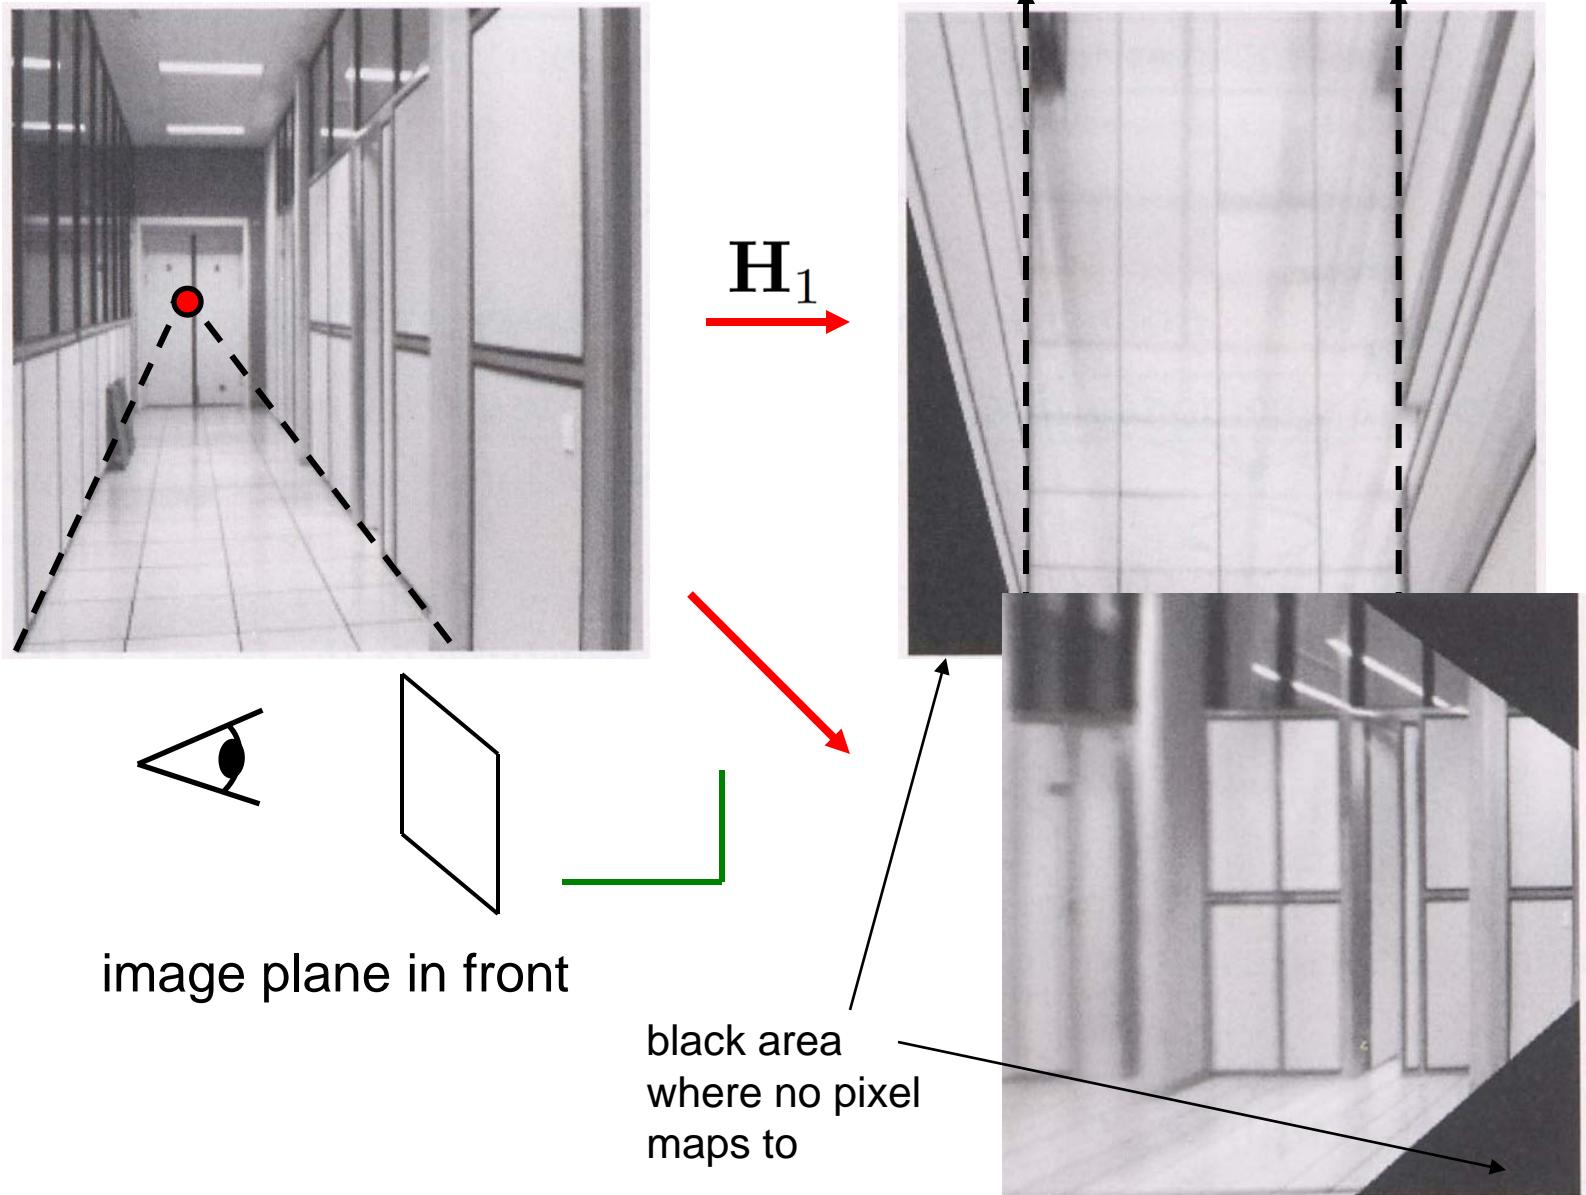


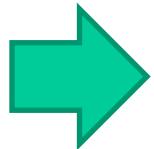
image plane below

# Image warping with homographies



# Homographies

---



Observation:

Rather than thinking of this as a 3D reprojection,  
think of it as a 2D **image warp** from one image to another

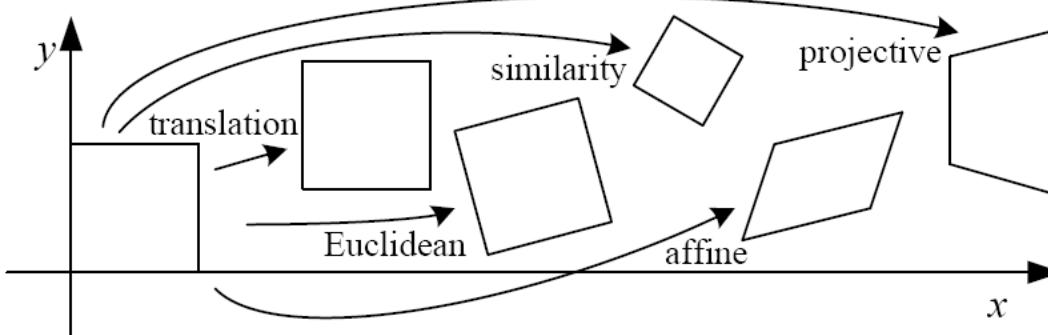
# Homographies

---

- Homographies ...
  - Affine transformations, and
  - Projective warps
- Properties of projective transformations:
  - Origin does not necessarily map to origin
  - Lines map to lines
  - Parallel lines do not necessarily remain parallel
  - Ratios are not preserved
  - Closed under composition

# 2D image transformations

---



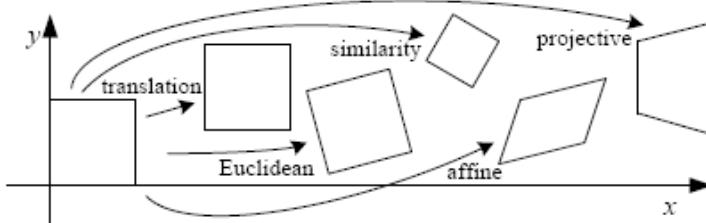
Name	Matrix	# D.O.F.	Preserves:	Icon
translation	$\begin{bmatrix} \mathbf{I} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	2	orientation + ...	
rigid (Euclidean)	$\begin{bmatrix} \mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	3	lengths + ...	
similarity	$\begin{bmatrix} s\mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	4	angles + ...	
affine	$\begin{bmatrix} \mathbf{A} \end{bmatrix}_{2 \times 3}$	6	parallelism + ...	
projective	$\begin{bmatrix} \tilde{\mathbf{H}} \end{bmatrix}_{3 \times 3}$	8	straight lines	

These transformations are a nested set of groups

- Closed under composition and inverse is a member

# 2D image transformations

Which transform is the right one for warping PP1 into PP2?  
e.g. translation, Euclidean, affine, projective



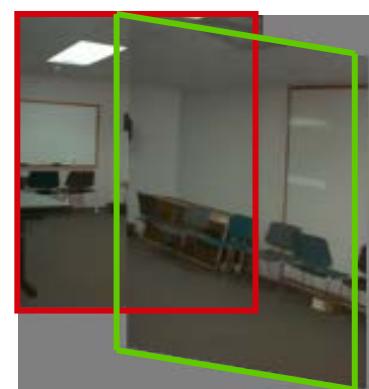
Translation

Affine

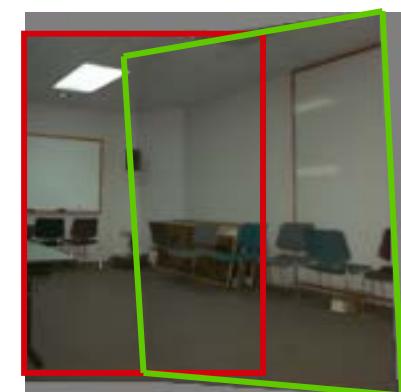
Perspective



2 unknowns



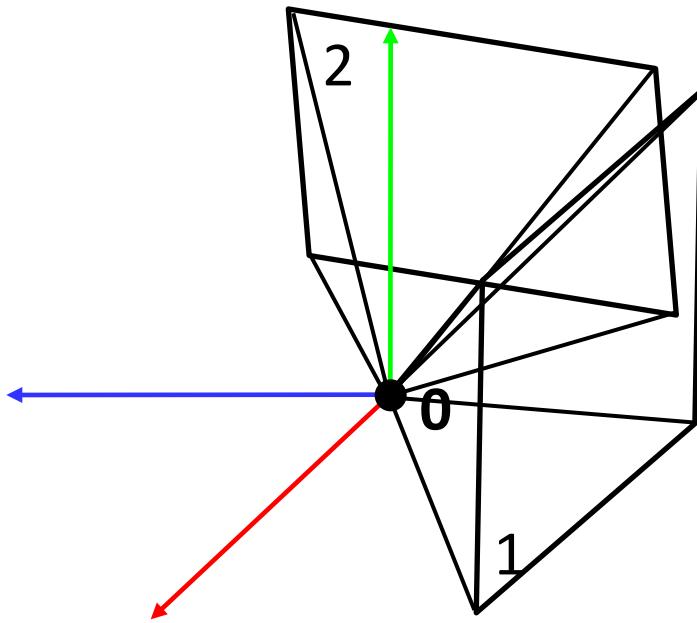
6 unknowns



8 unknowns

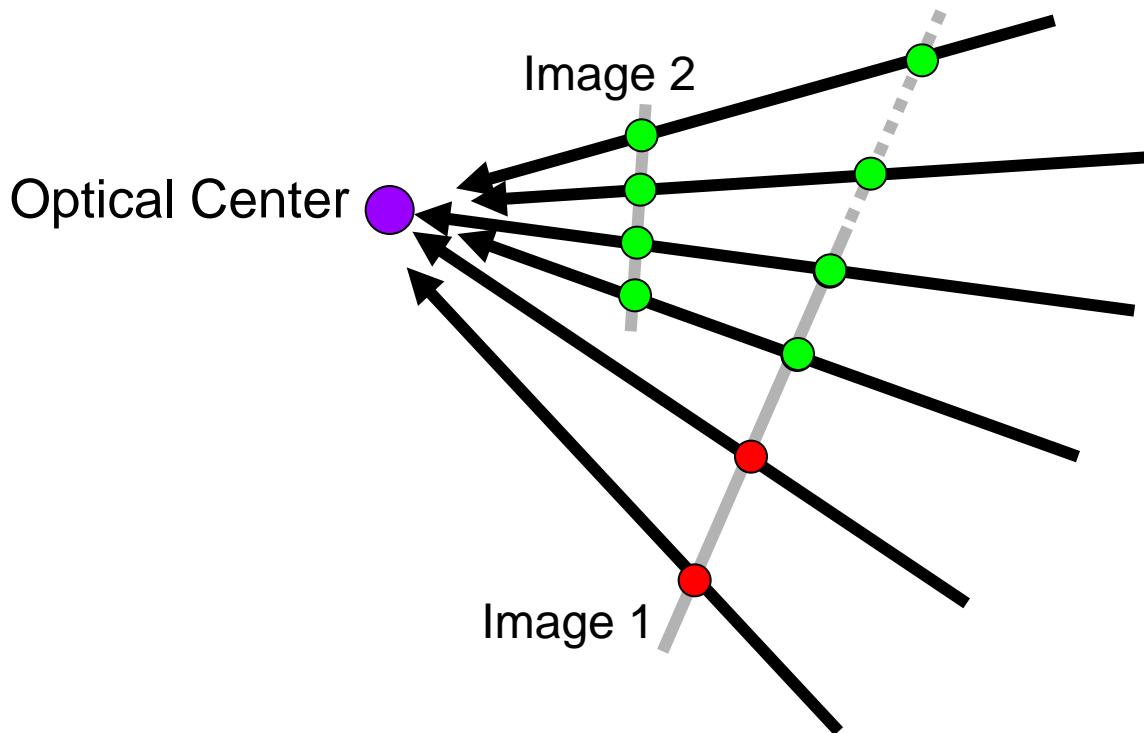
# Geometric setup of mosaics

---



# What is the transformation?

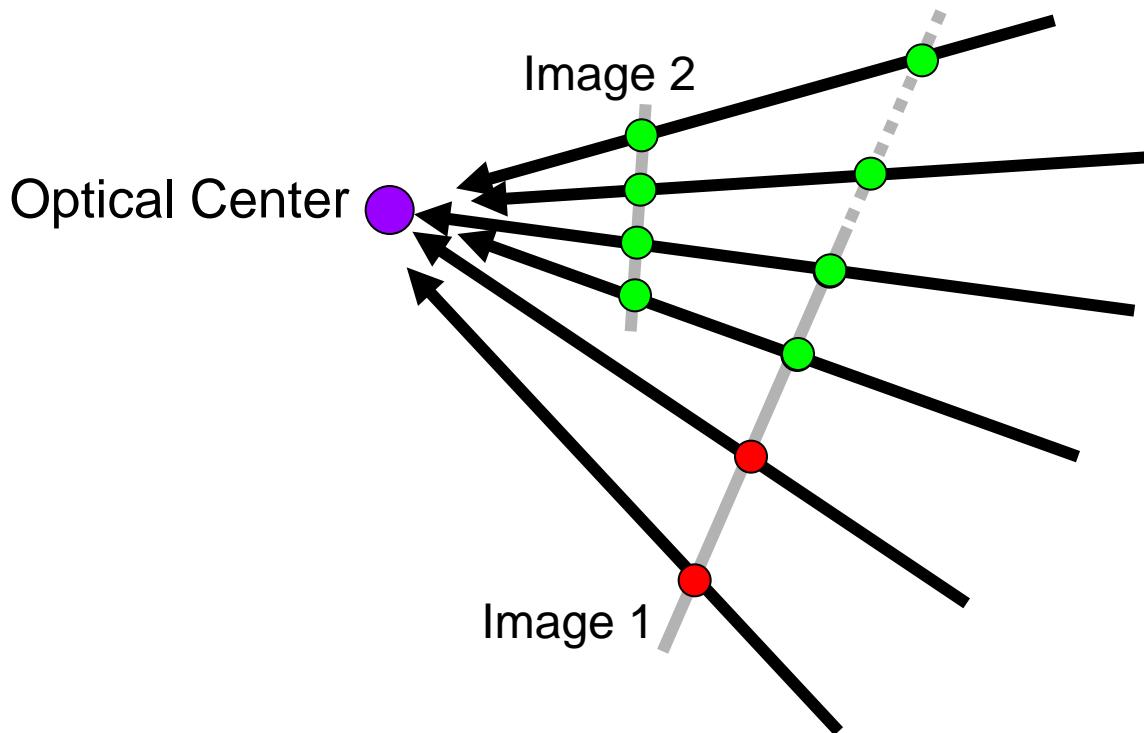
---



$$\begin{pmatrix} x_1 \\ y_1 \\ 1 \end{pmatrix} \sim \mathbf{K}_1 \begin{pmatrix} X \\ Y \\ Z \end{pmatrix}$$

# What is the transformation?

---

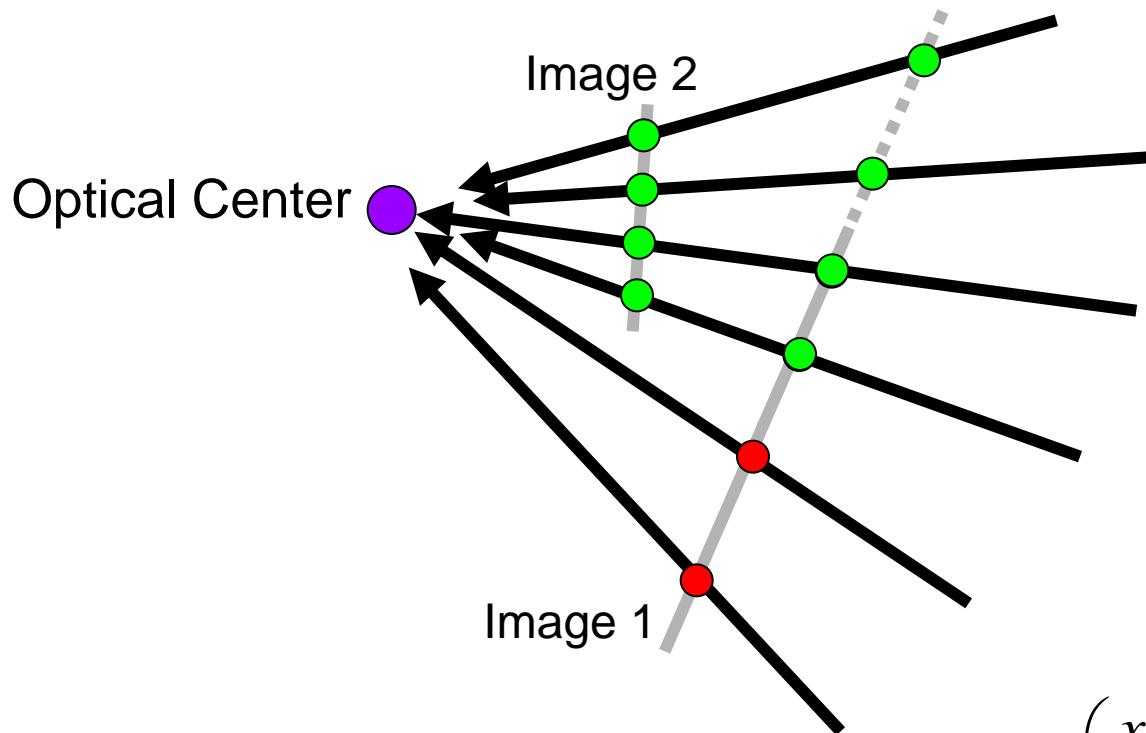


$$\begin{pmatrix} x_1 \\ y_1 \\ 1 \end{pmatrix} \sim \mathbf{K}_1 \begin{pmatrix} X \\ Y \\ Z \end{pmatrix}$$

$$\begin{pmatrix} x_2 \\ y_2 \\ 1 \end{pmatrix} \sim \mathbf{K}_2 \mathbf{R} \begin{pmatrix} X \\ Y \\ Z \end{pmatrix}$$

# What is the transformation?

---



$$\begin{pmatrix} x_1 \\ y_1 \\ 1 \end{pmatrix} \sim \mathbf{K}_1 \begin{pmatrix} X \\ Y \\ Z \end{pmatrix}$$

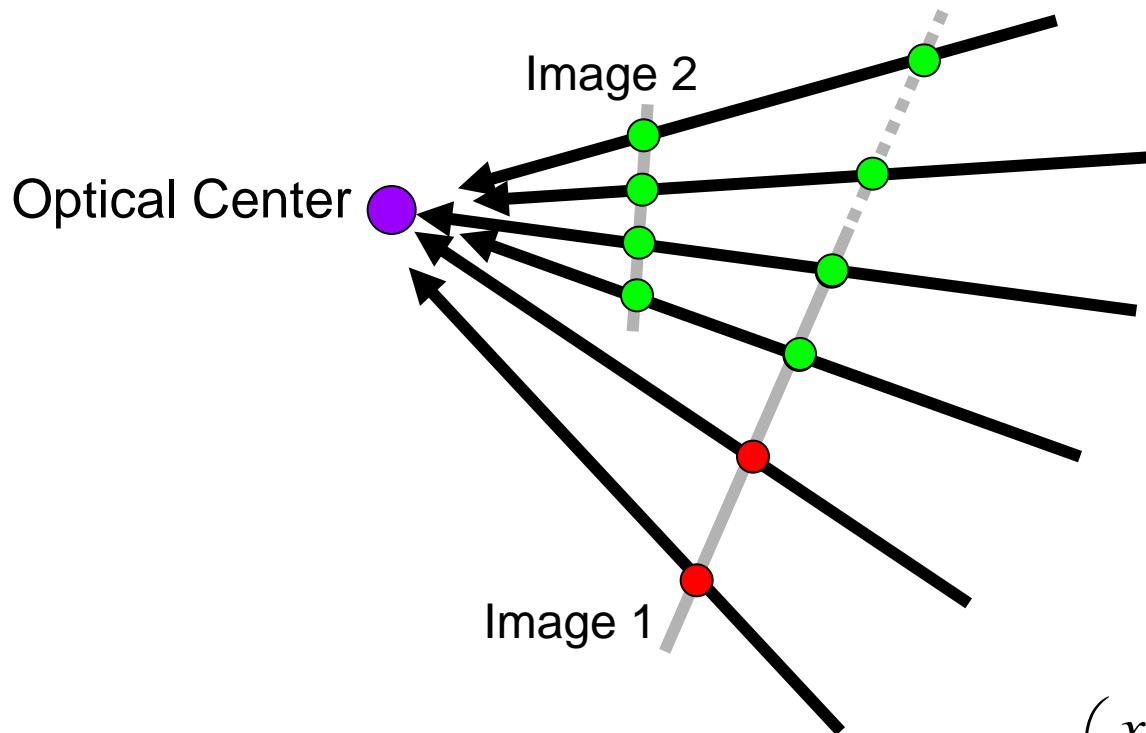
$$\begin{pmatrix} x_2 \\ y_2 \\ 1 \end{pmatrix} \sim \mathbf{K}_2 \mathbf{R} \begin{pmatrix} X \\ Y \\ Z \end{pmatrix}$$

$$\begin{pmatrix} x_2 \\ y_2 \\ 1 \end{pmatrix} \sim \mathbf{K}_2 \mathbf{R} \mathbf{K}_1^{-1} \begin{pmatrix} x_1 \\ y_1 \\ 1 \end{pmatrix}$$

3x3 matrix  
also called Homography

# What is the transformation?

---



$$\begin{pmatrix} x_1 \\ y_1 \\ 1 \end{pmatrix} \sim \mathbf{K}_1 \begin{pmatrix} X \\ Y \\ Z \end{pmatrix}$$

$$\begin{pmatrix} x_2 \\ y_2 \\ 1 \end{pmatrix} \sim \mathbf{K}_2 \mathbf{R} \begin{pmatrix} X \\ Y \\ Z \end{pmatrix}$$

$$\begin{pmatrix} x_2 \\ y_2 \\ 1 \end{pmatrix} \sim \mathbf{K}_2 \mathbf{R} \mathbf{K}_1^{-1} \begin{pmatrix} x_1 \\ y_1 \\ 1 \end{pmatrix}$$

3x3 matrix

also called Homography

Project from image to 3D ray

Rotate the ray by camera motion

Project back into new (source) image

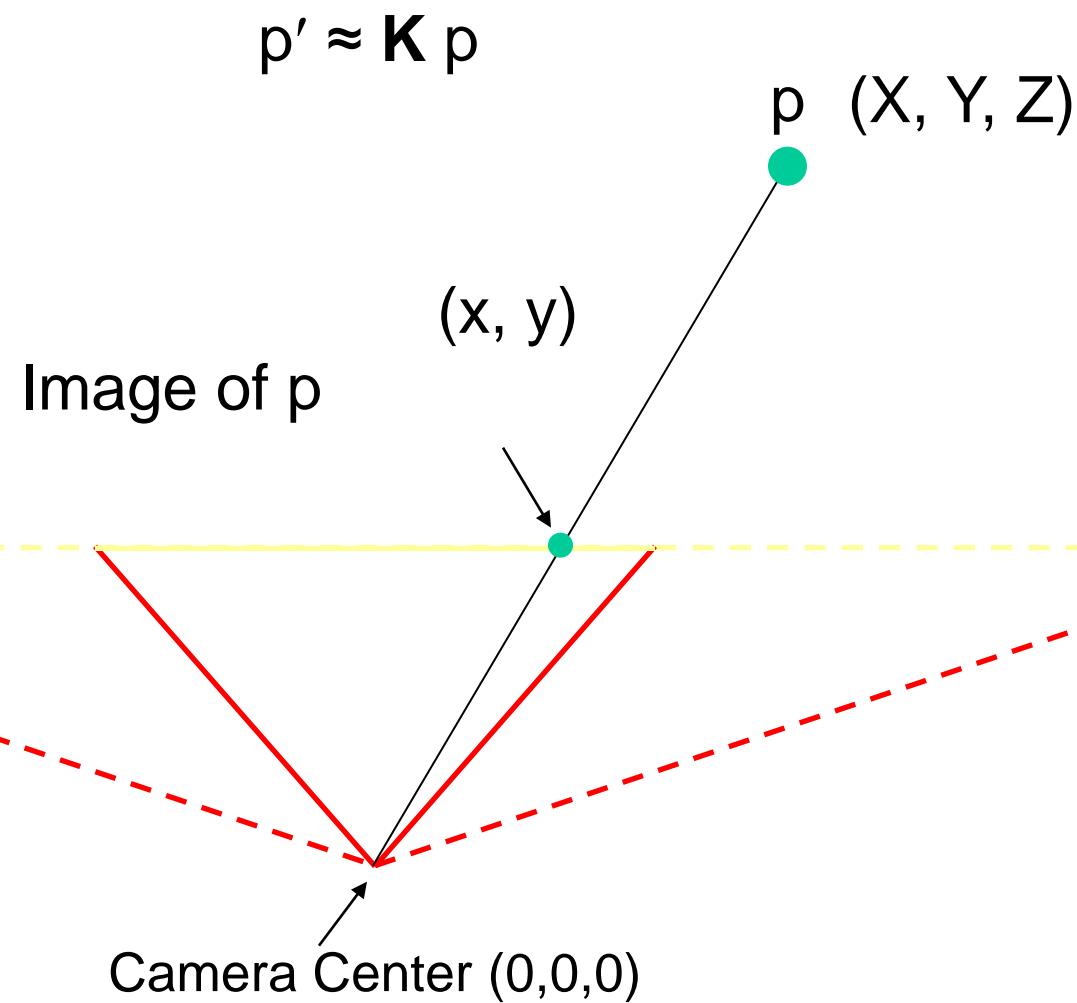
# Combining Images into Panoramas

---

- Theorem: Any 2 images of an arbitrary scene taken with *same* camera center are related by  $p_2 \approx \mathbf{K} \mathbf{R} \mathbf{K}^{-1} p_1$  where  $p_1$  and  $p_2$  are homogeneous coords of 2 corresponding points,  $\mathbf{K}$  is  $3 \times 3$  camera calibration matrix, and  $\mathbf{R}$  is  $3 \times 3$  rotation matrix
- $\mathbf{K} \mathbf{R} \mathbf{K}^{-1}$  is  $3 \times 3$  matrix called the “homography induced by the plane at infinity”

# Another example

---

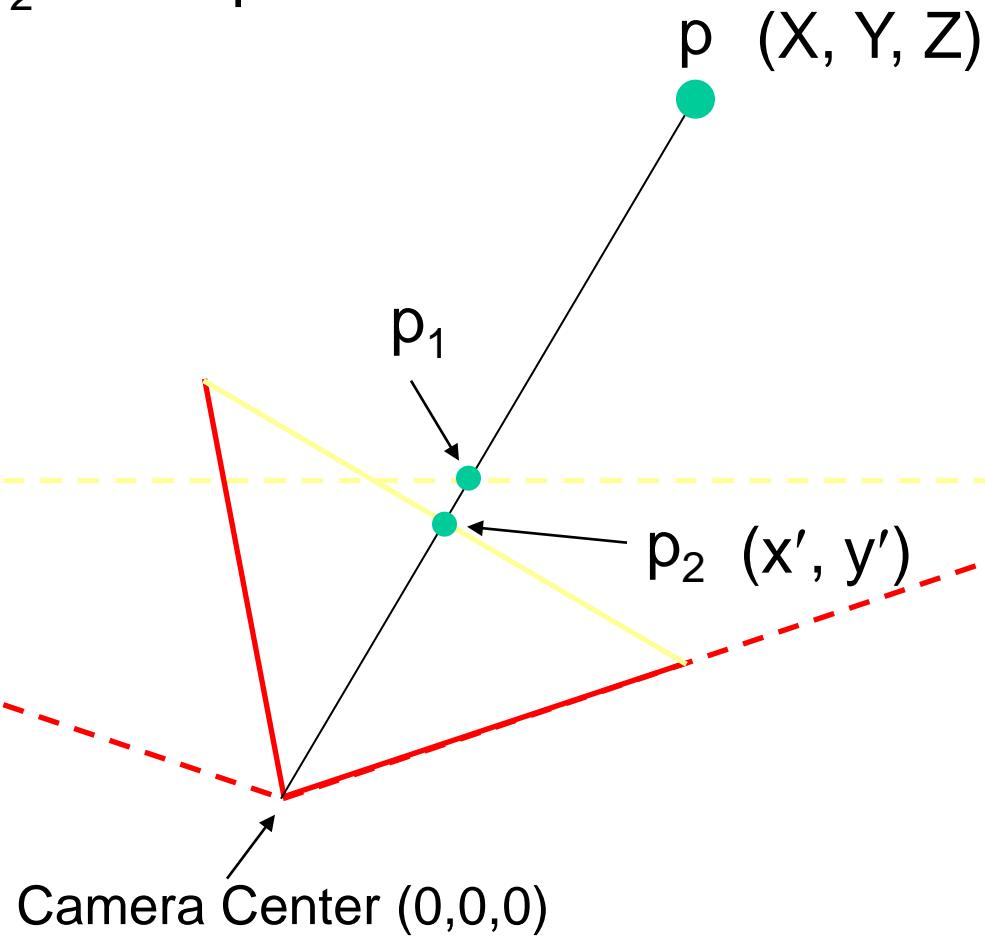


# Another example

---

$$p_1 \approx K p$$

$$p_2 \approx K R p$$



# Another example

---

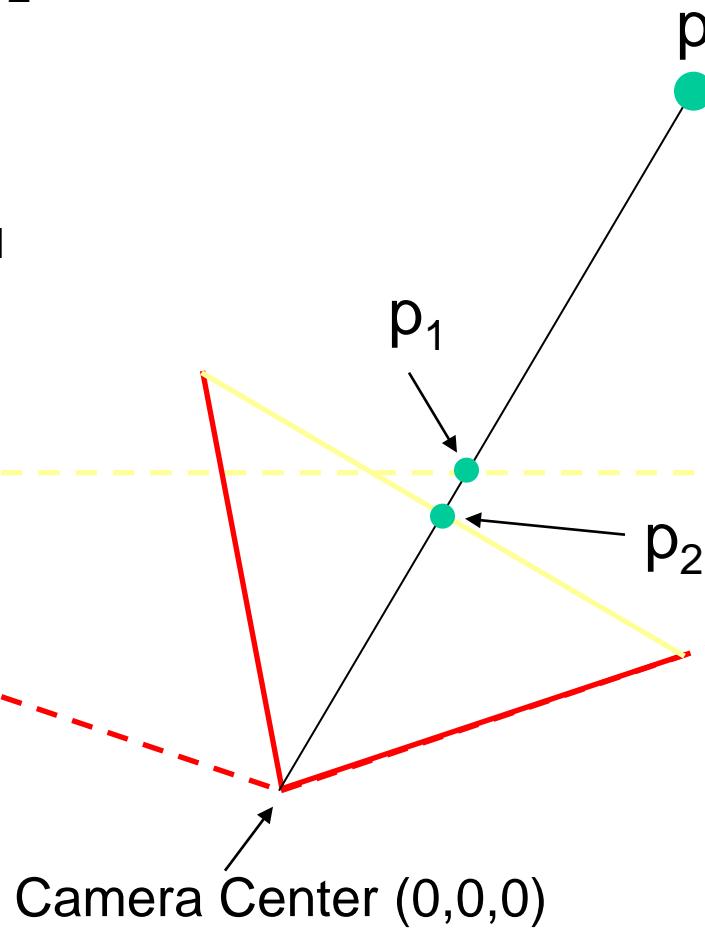
$$p_1 \approx \mathbf{K} p$$

$$p_2 \approx \mathbf{K} R p$$

$$\mathbf{K}^{-1} p_1 \approx p$$

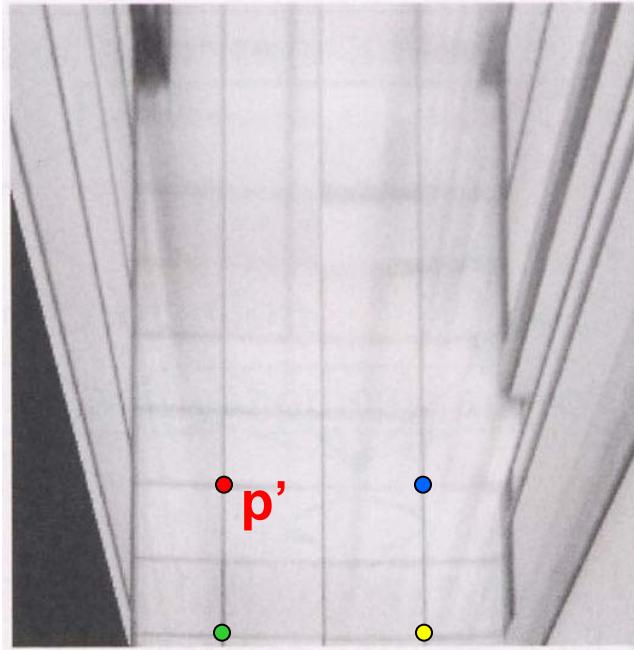
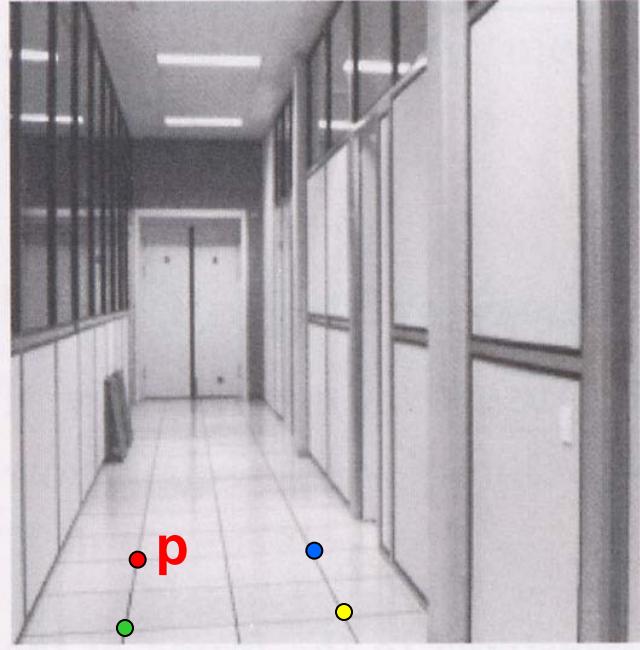
$$p_2 \approx \boxed{\mathbf{K} R \mathbf{K}^{-1}} p_1$$

homography



# Image rectification

---



To un warp (rectify) an image

- Find the homography  $\mathbf{H}$  given a set of  $\mathbf{p}$  and  $\mathbf{p}'$  pairs
- How many correspondences are needed?

# Solving for homographies

---

$$\begin{aligned} \mathbf{p}' &= \mathbf{H}\mathbf{p} \\ \begin{bmatrix} wx' \\ wy' \\ w \end{bmatrix} &= \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} \end{aligned}$$

Can set scale factor  $i=1$  and  $w=1$ . So, there are 8 unknowns.

Set up a system of linear equations:

$$\mathbf{A}\mathbf{h} = \mathbf{b}$$

where vector of unknowns  $\mathbf{h} = [a,b,c,d,e,f,g,h]^T$

Need at least 8 eqs, but the more the better...

Solve for  $\mathbf{h}$ . If overconstrained, solve using least-squares:

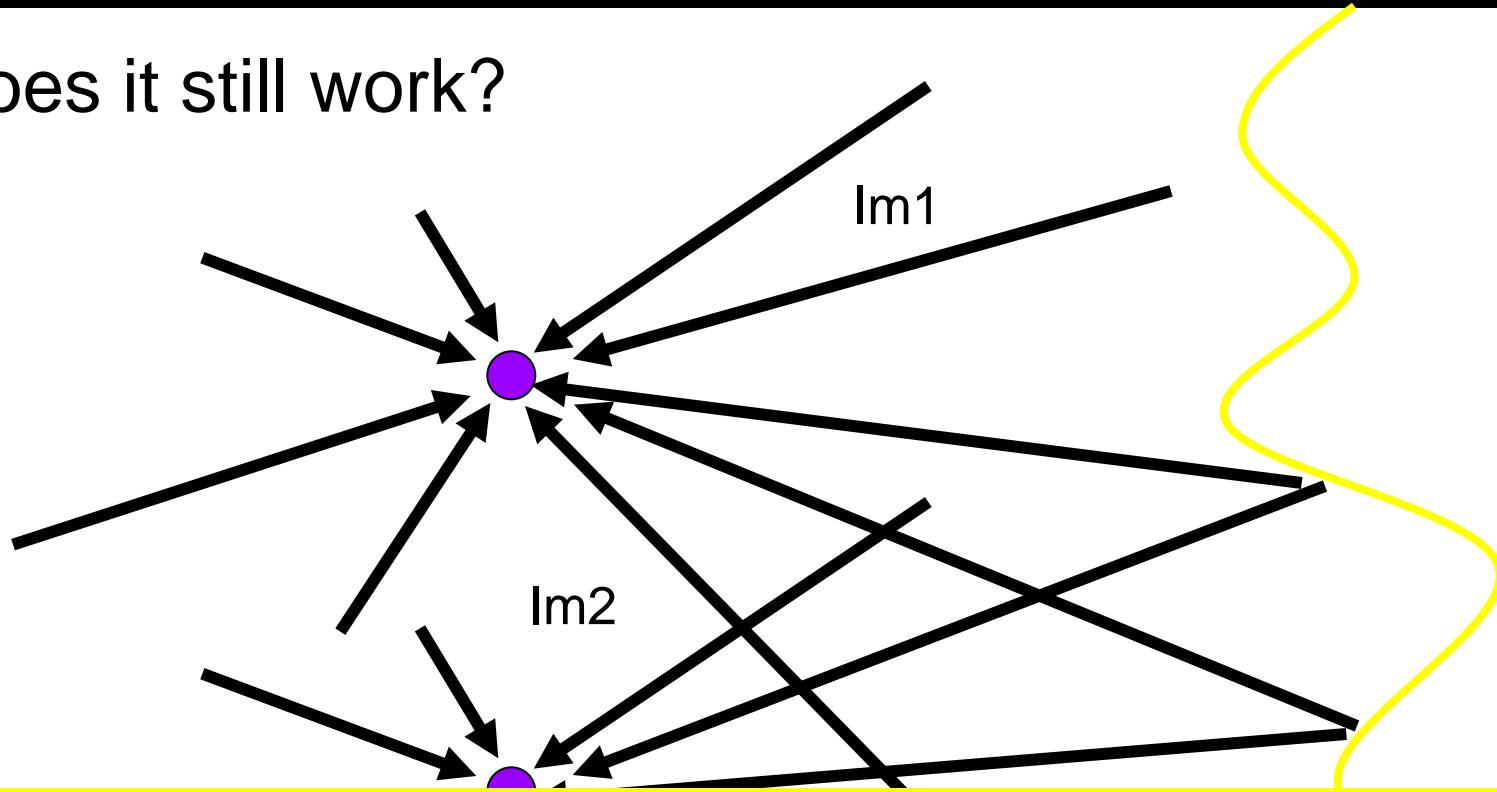
Can be done in Matlab

$$\min \|A\mathbf{h} - \mathbf{b}\|^2$$

# changing camera center

---

- Does it still work?



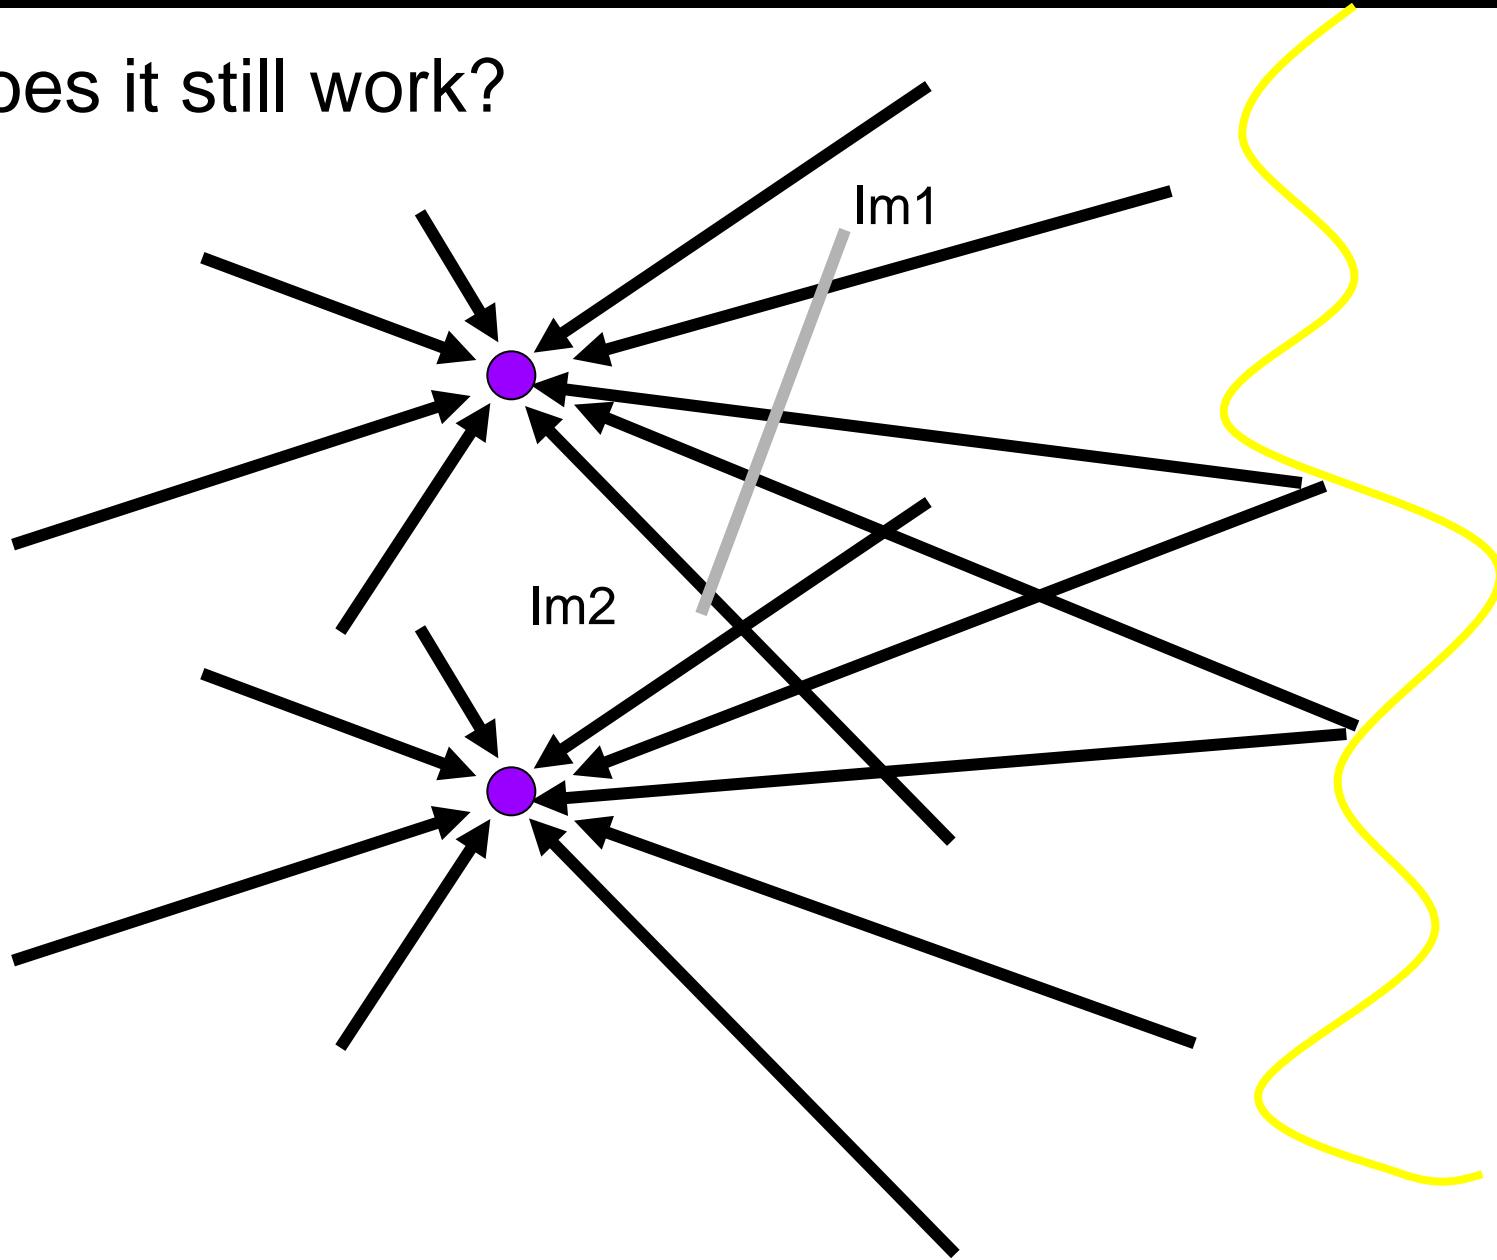
The mosaic has a natural interpretation in 3D

- The images are reprojected onto a common plane
- The mosaic is formed on this plane
- Mosaic is a *synthetic wide-angle camera*

# changing camera center

---

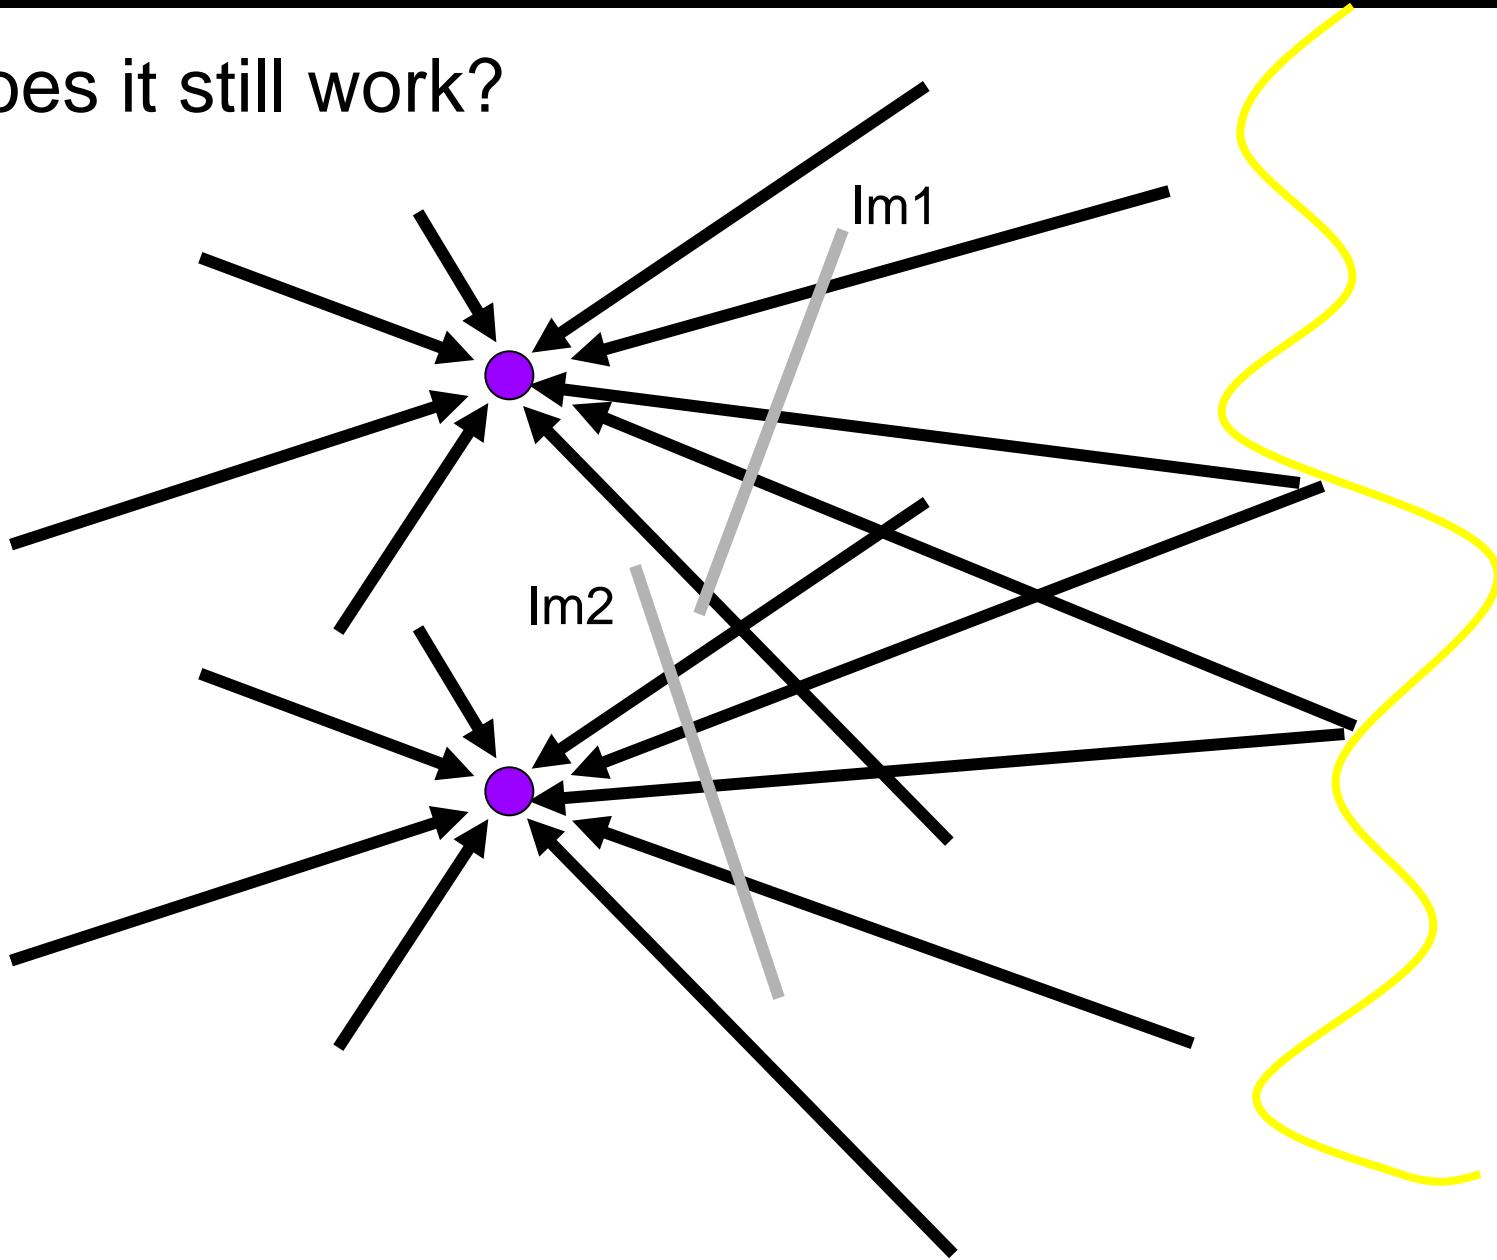
- Does it still work?



# changing camera center

---

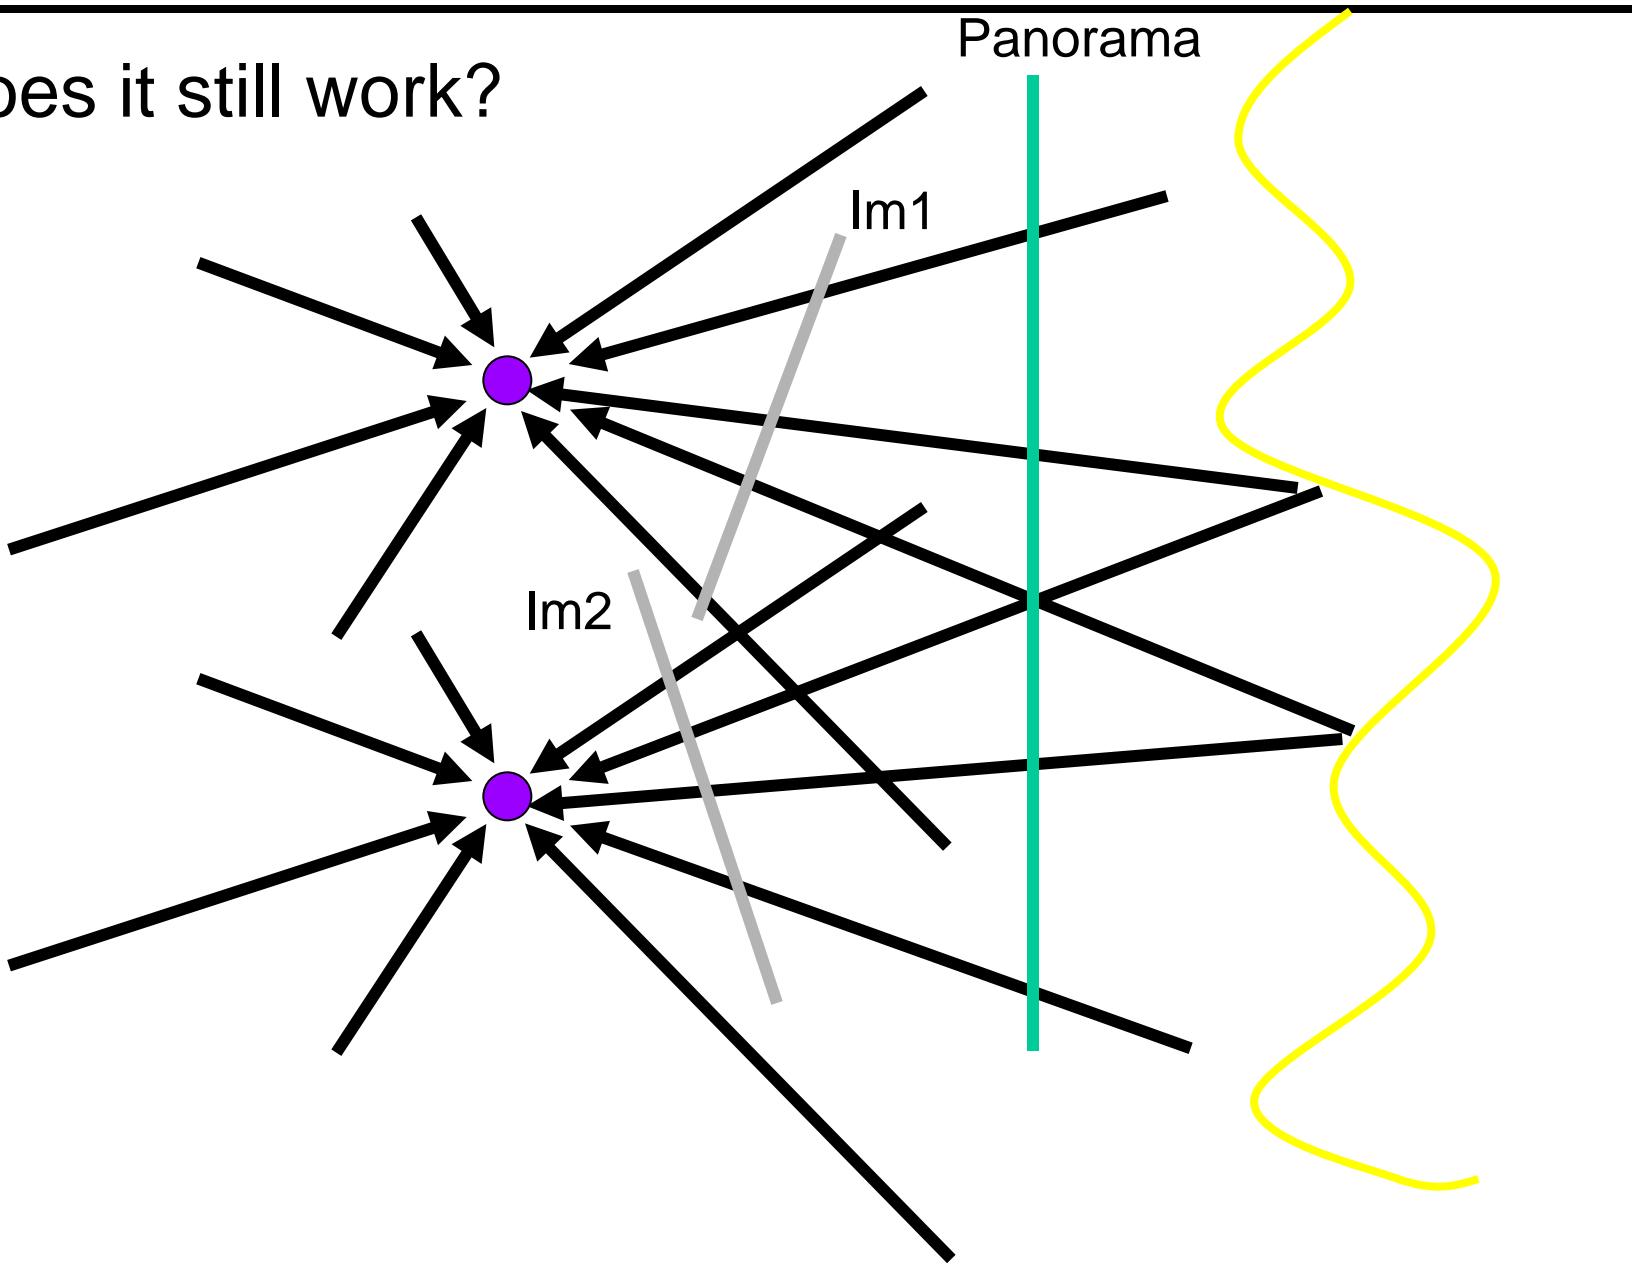
- Does it still work?



# changing camera center

---

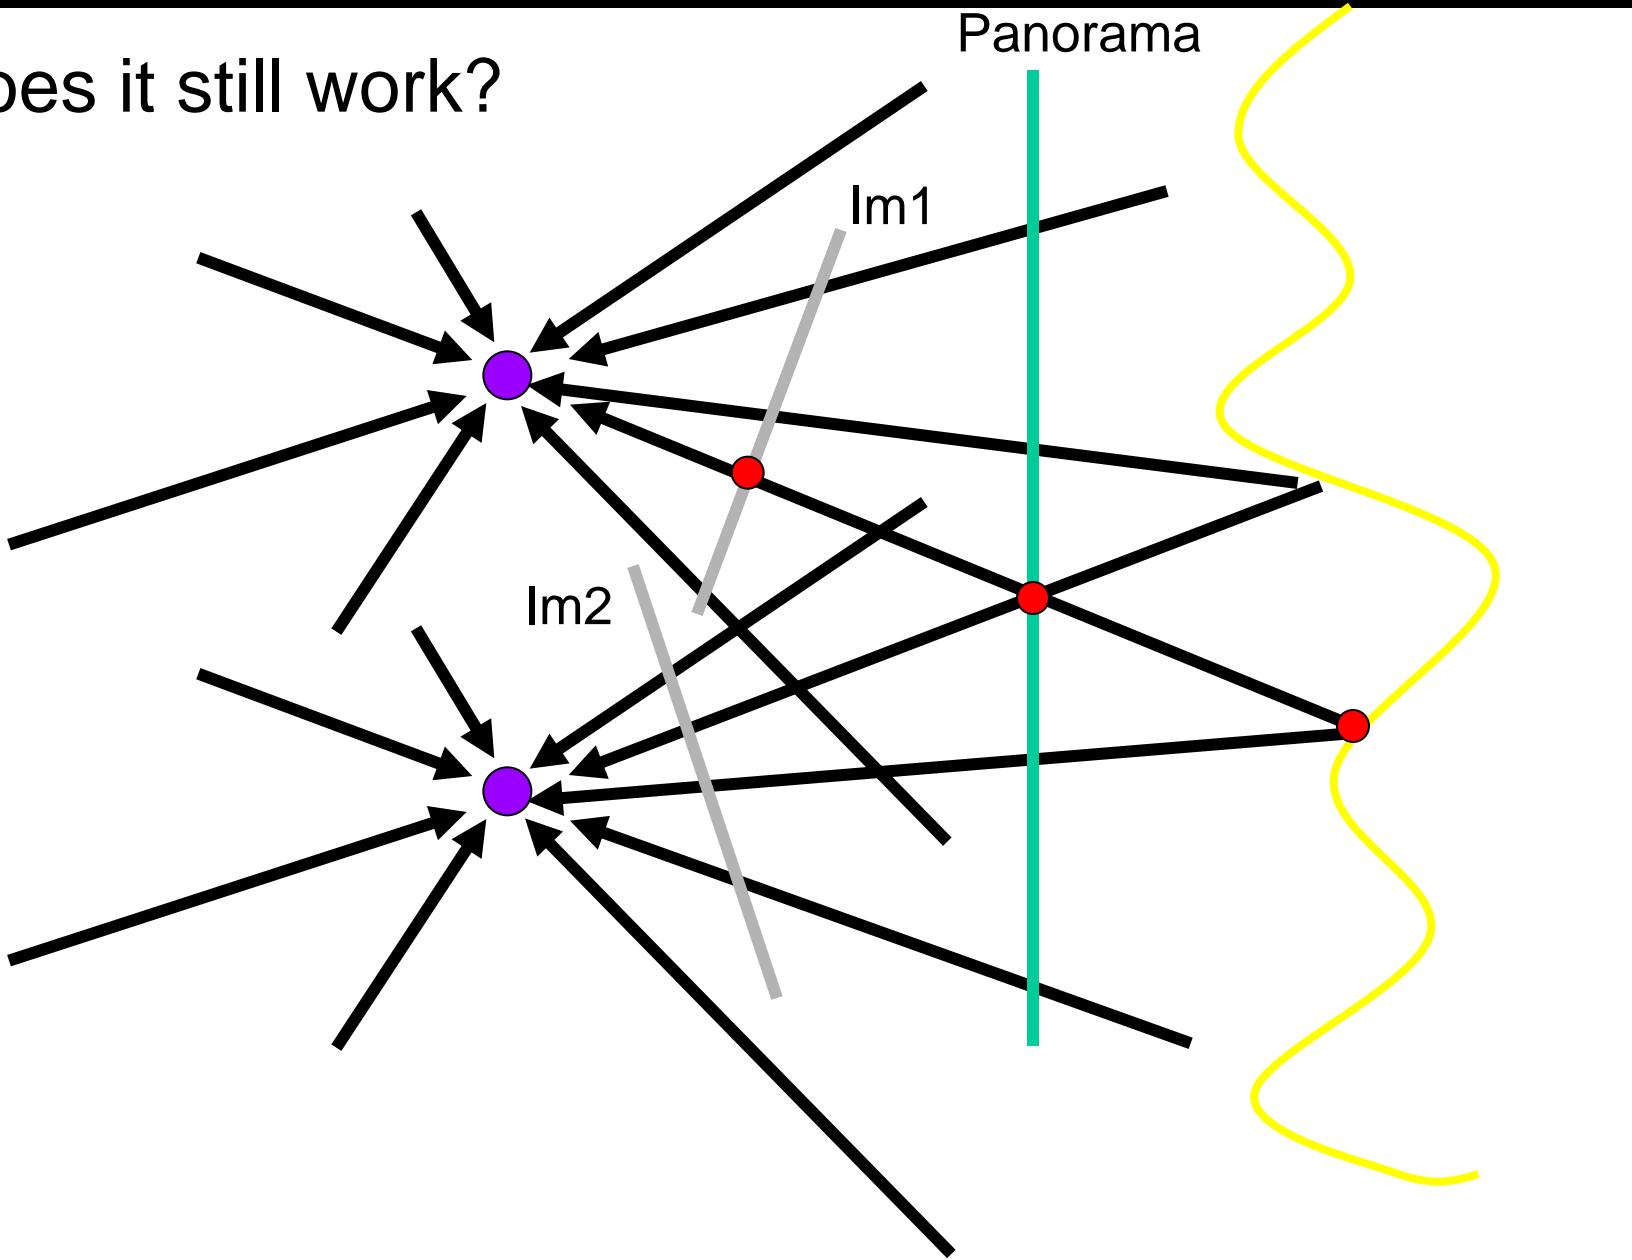
- Does it still work?



# changing camera center

---

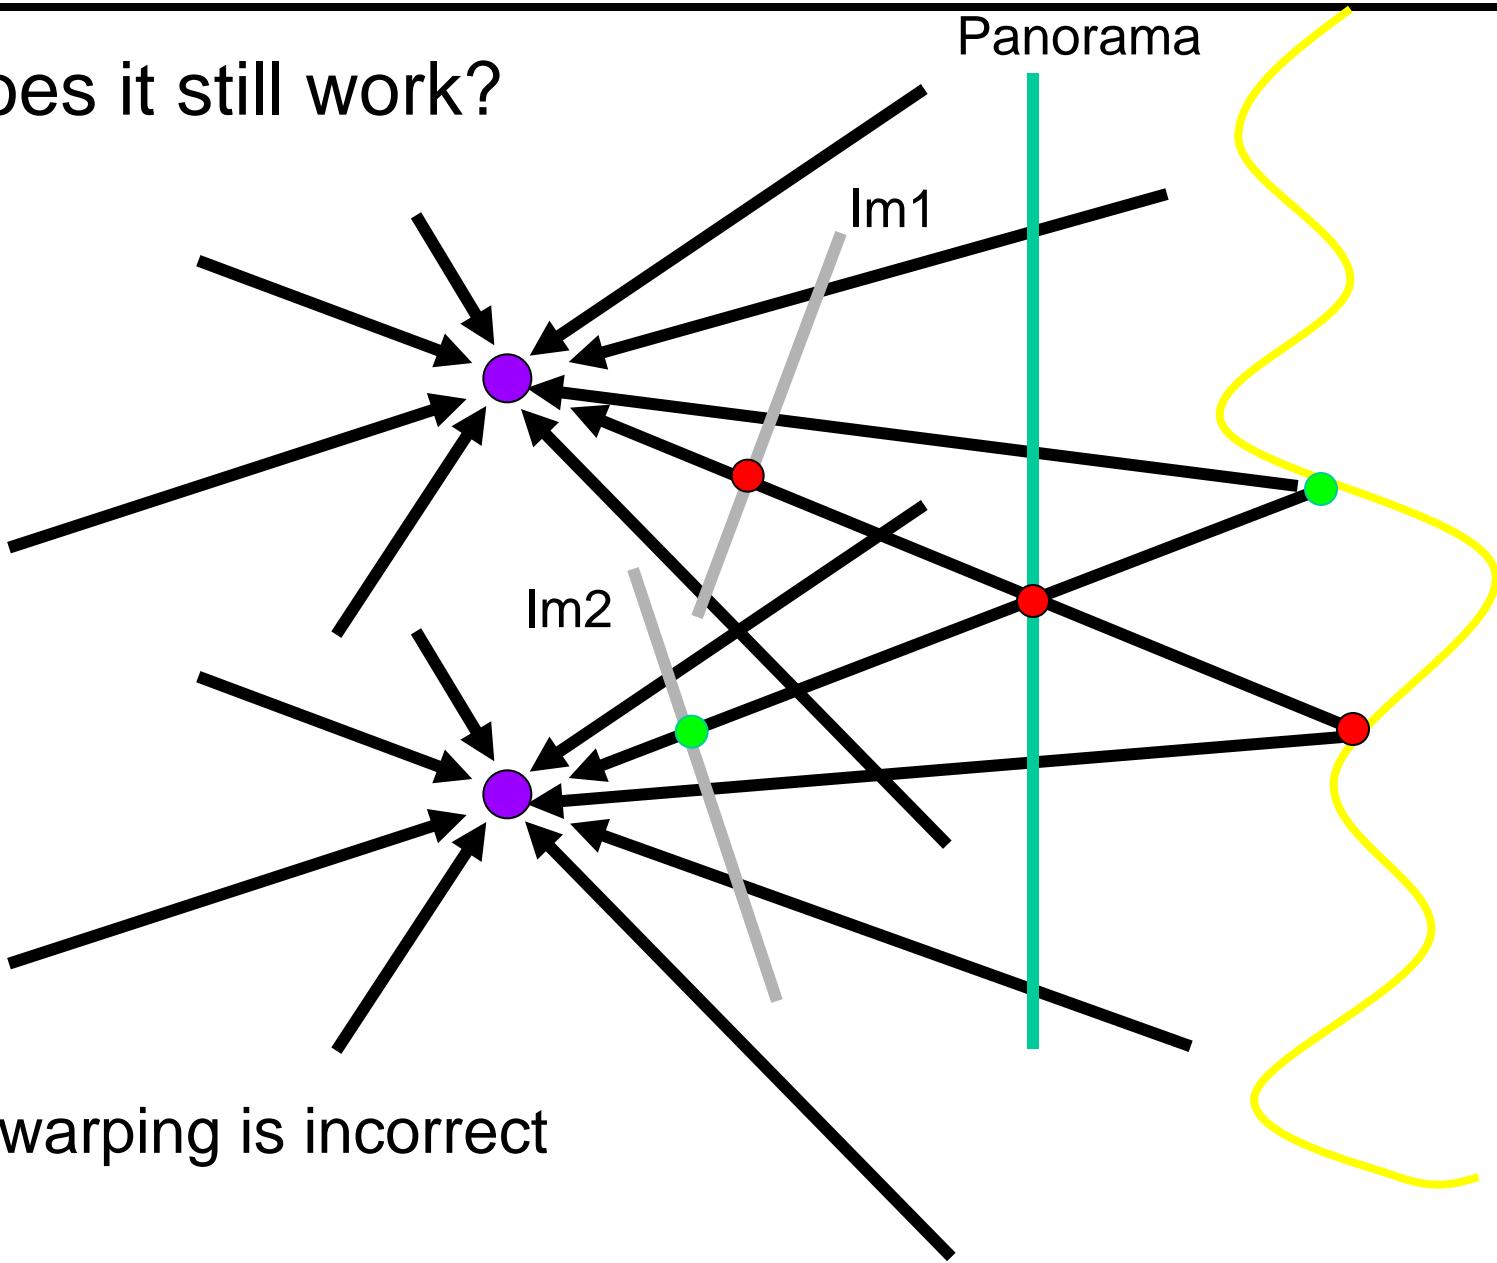
- Does it still work?



# changing camera center

---

- Does it still work?

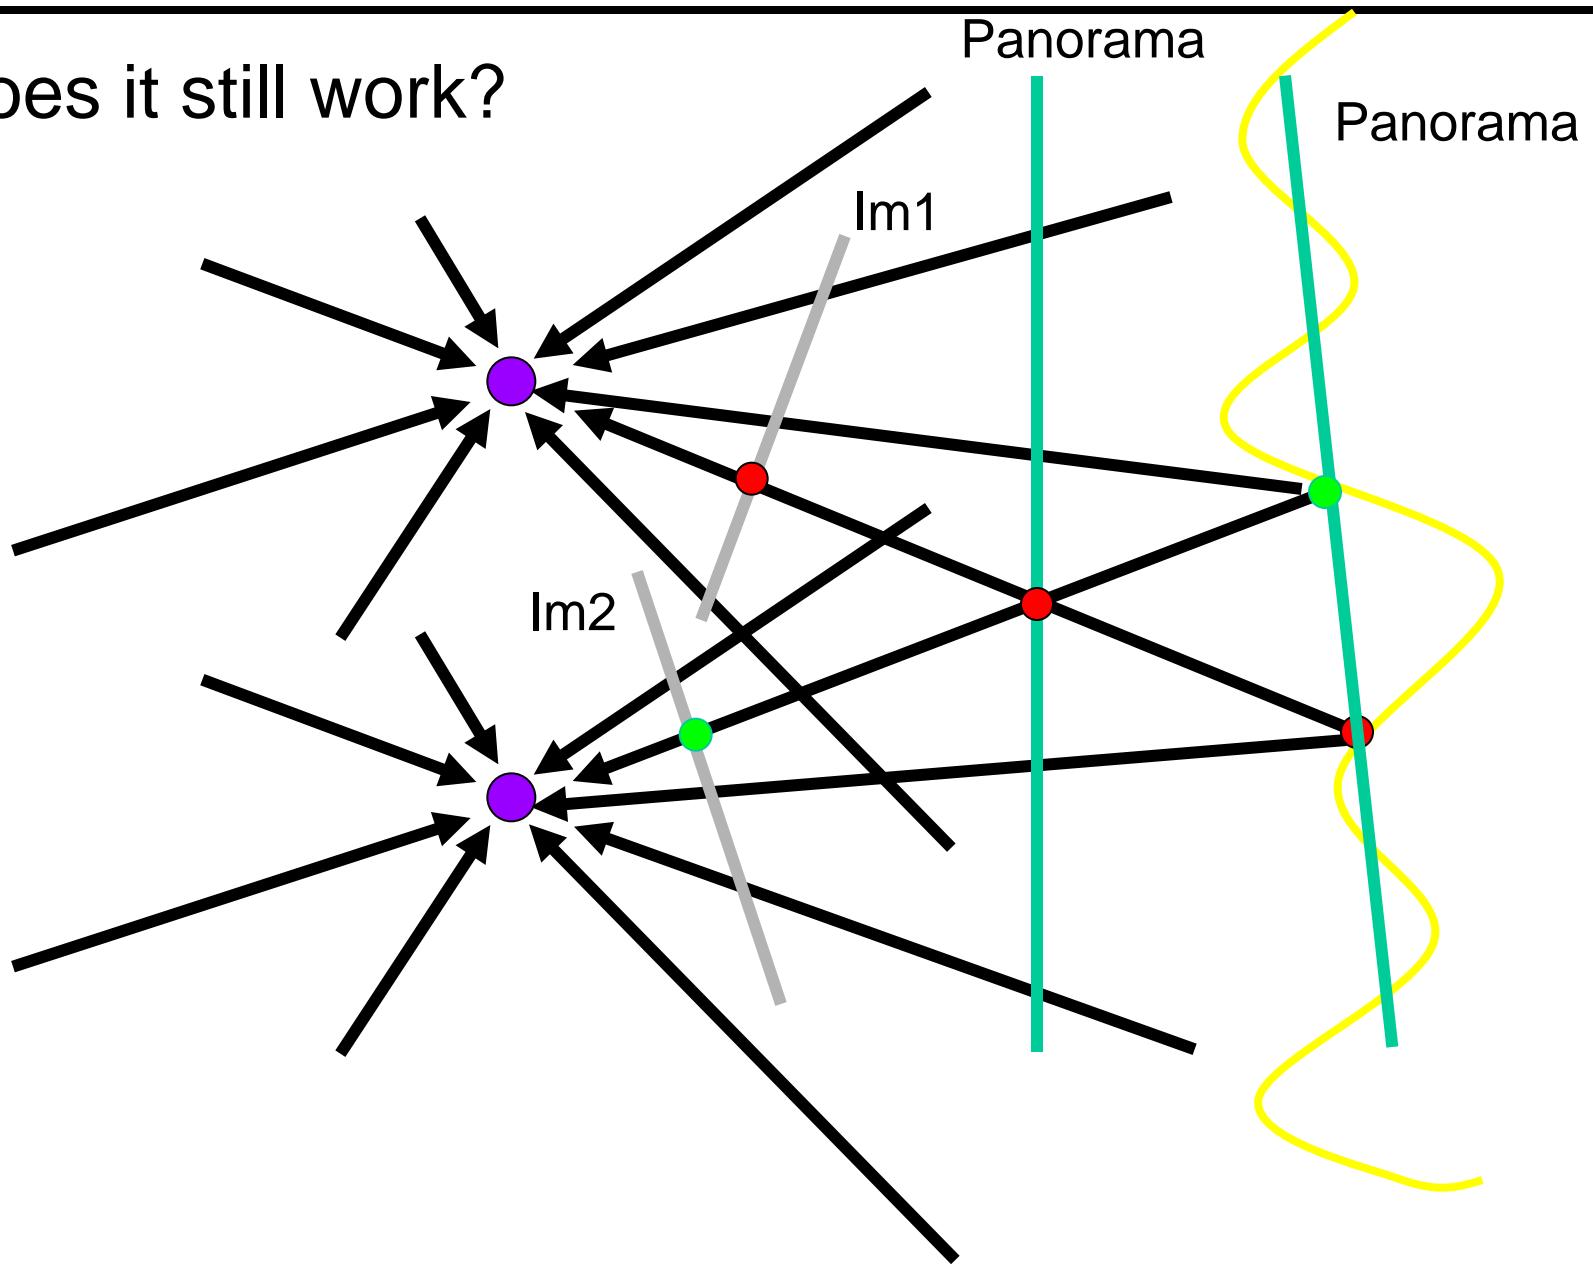


Your warping is incorrect

# changing camera center

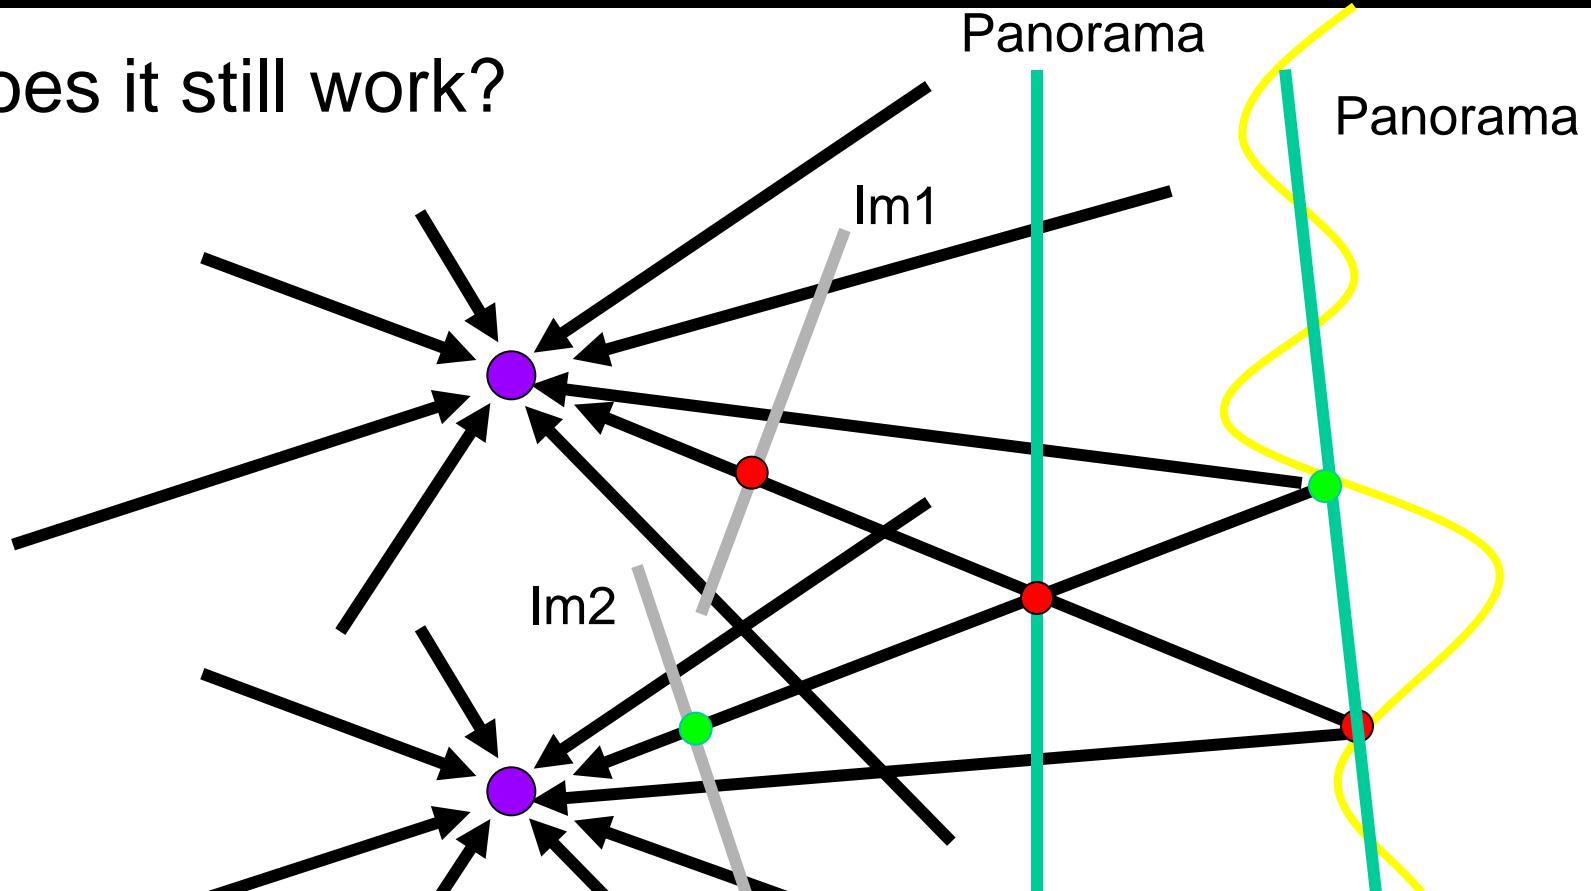
---

- Does it still work?



# changing camera center

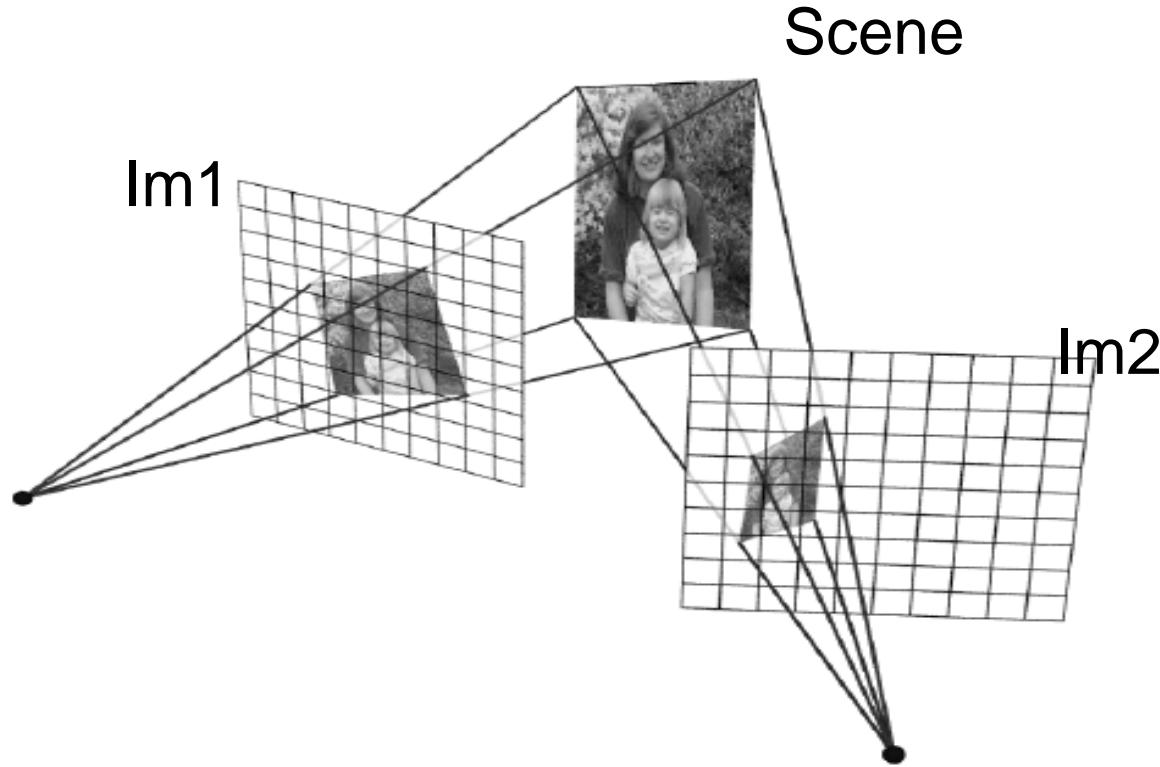
- Does it still work?



So either far away or no depth/relief within the scene

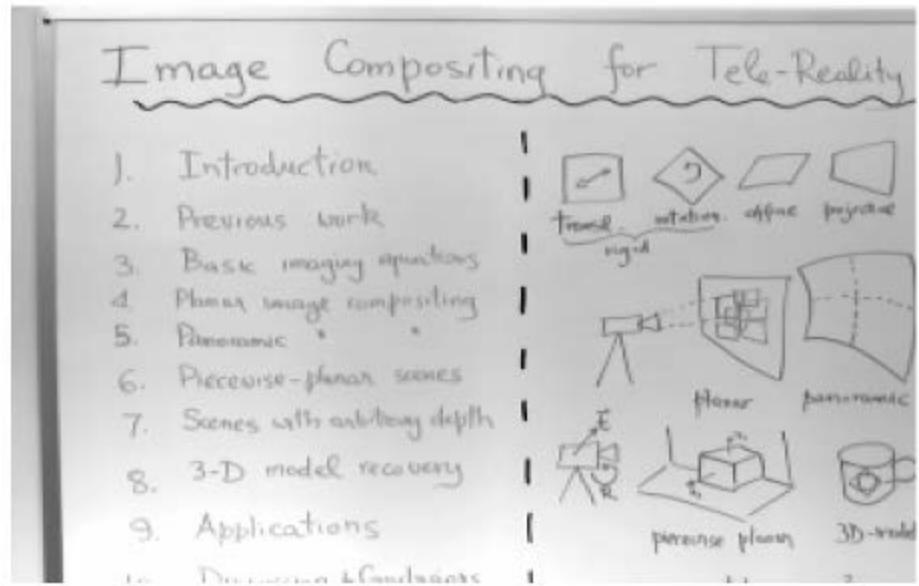
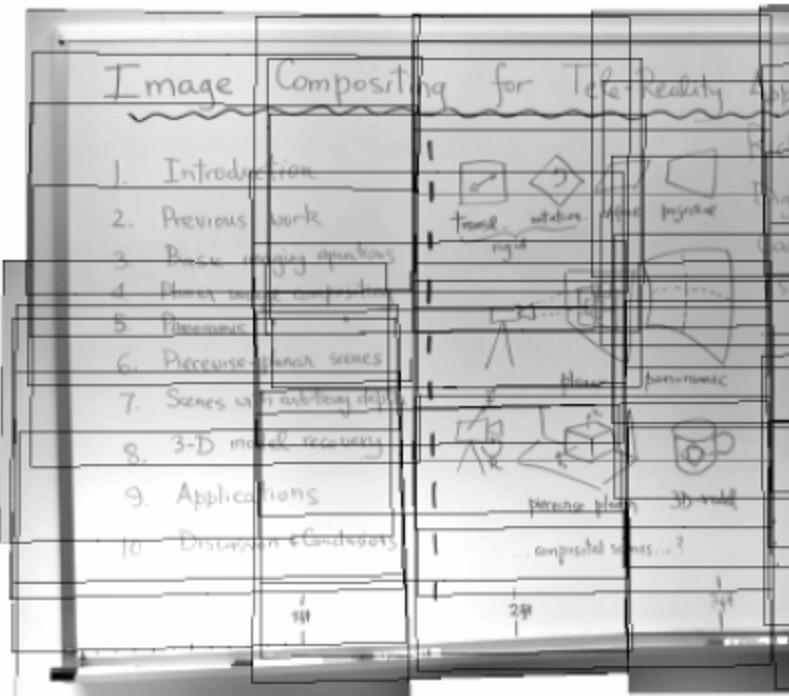
# Planar scene (far away)

---

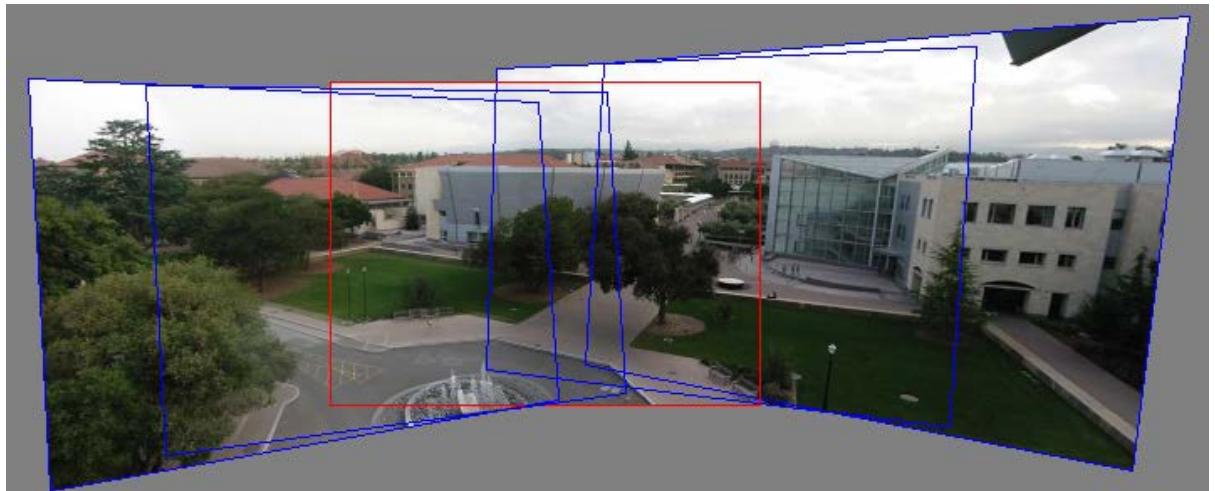
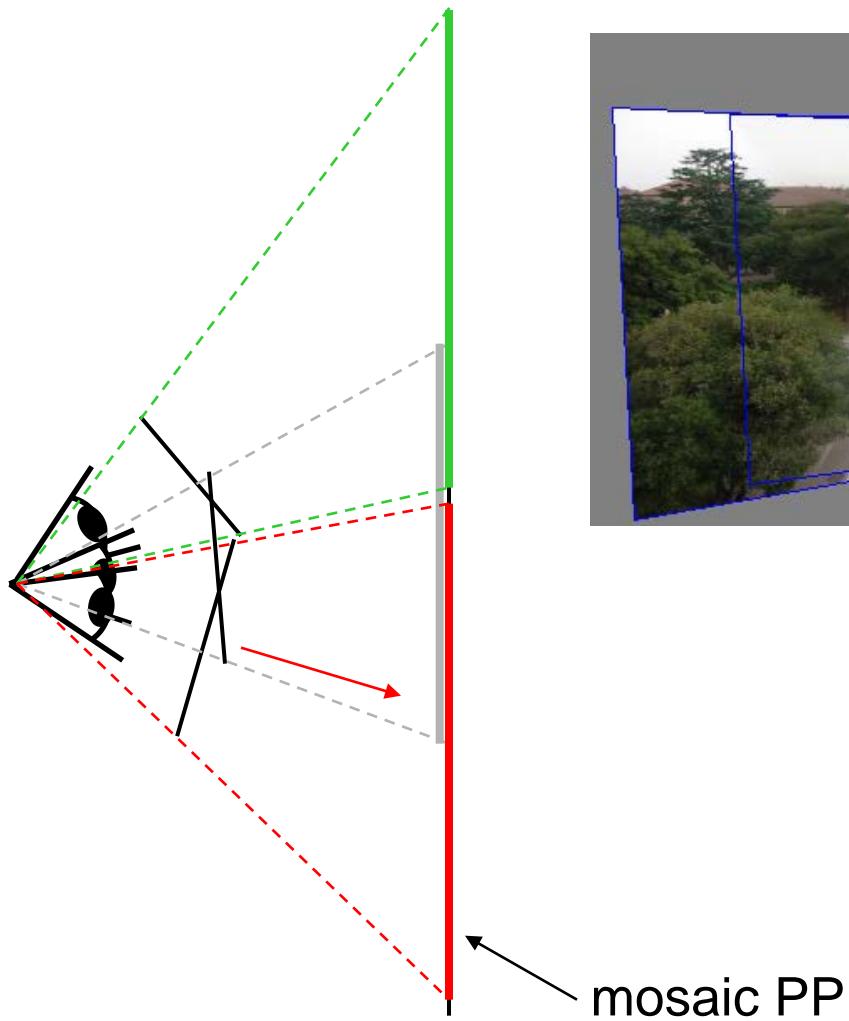


- If scene is planar, we are OK!
- This is how big aerial photographs are made

# Planar mosaic Examples

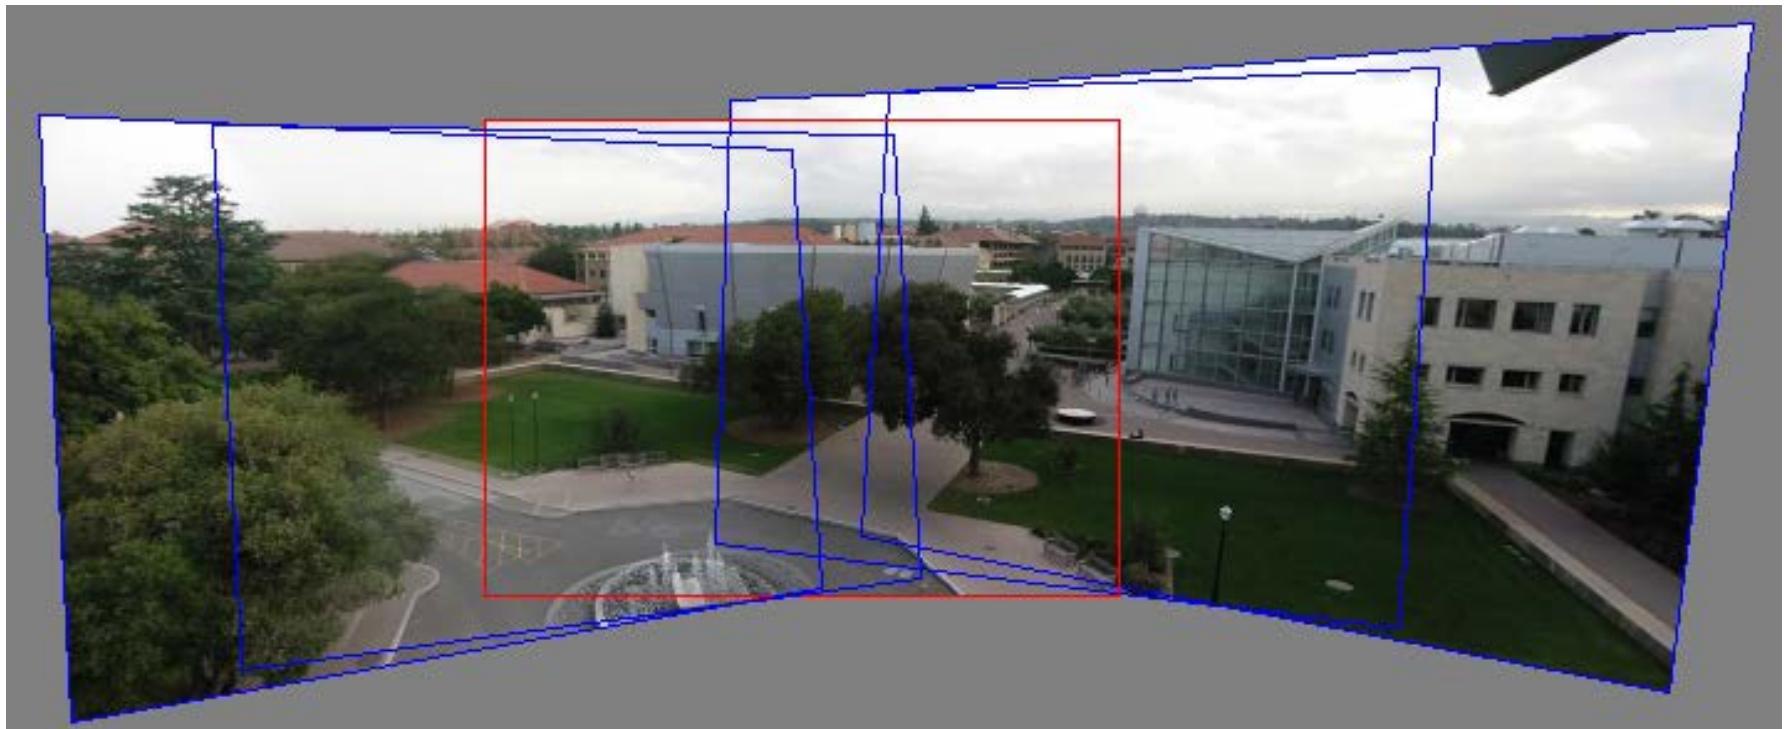


# Projecting images onto a common plane



# Can use homography to create a 360 panorama

---



1. Pick one image (red)
2. Warp the other images towards it (usually, one by one)
3. blend

# Recap:

- With enough images from the same optical center, we can create panorama.
- If the camera moves, we can't in general
- If the scene is planar or faraway, we are OK.

# Moving on...

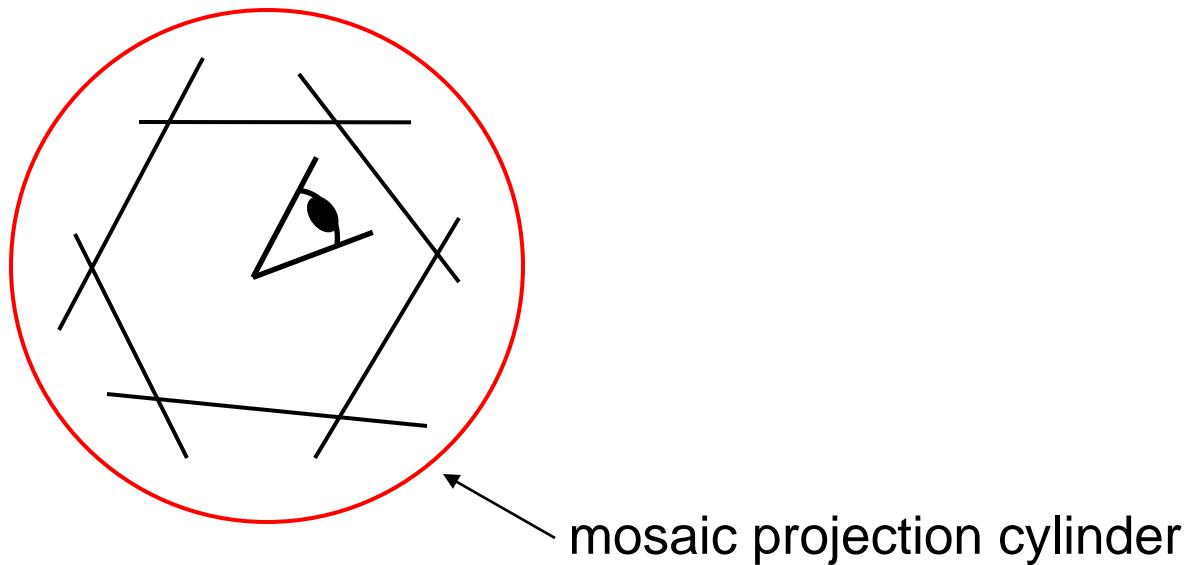
---



# Should use Cylindrical Projection

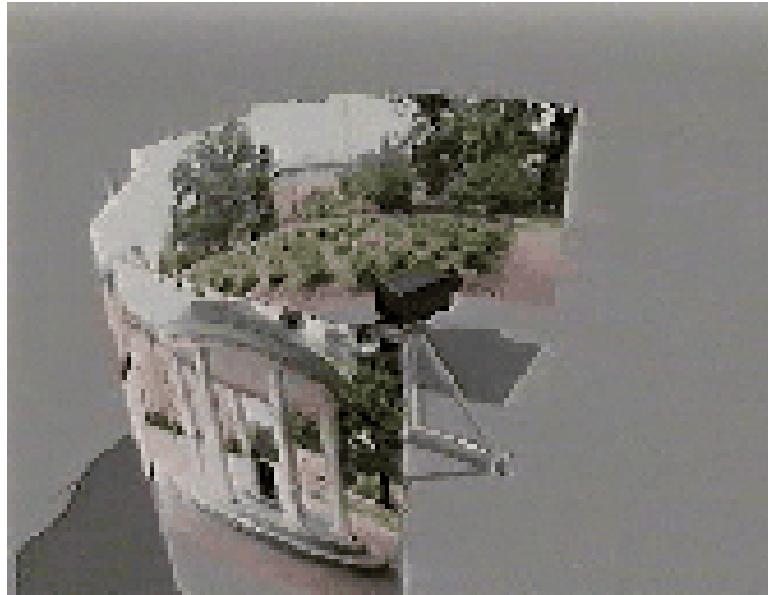
---

When we cannot choose a “reference” image (the red one) to map everything else to...



# Cylindrical panoramas:

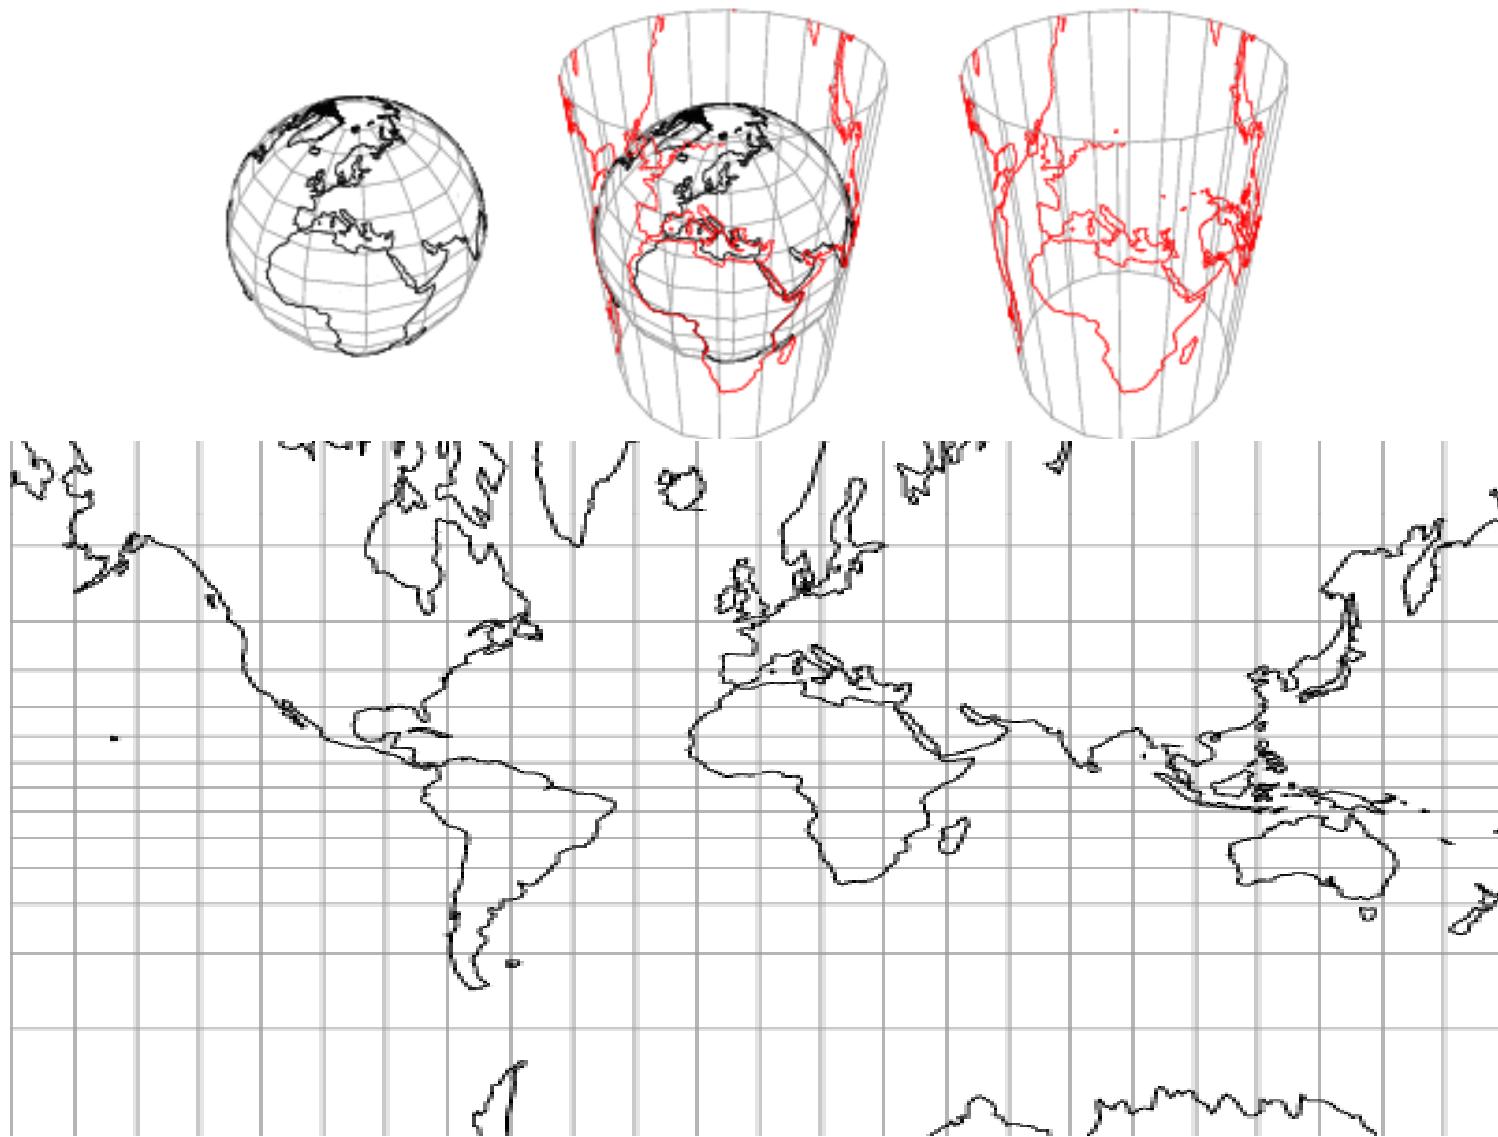
---



- Steps
  - Reproject each image onto a cylinder
  - Align and Blend
  - Output the resulting mosaic

# Cylindrical projection (An Example)

---



# Taking pictures

---



Kaidan panoramic tripod head  
Assume principal axis is parallel  
to horizontal plane

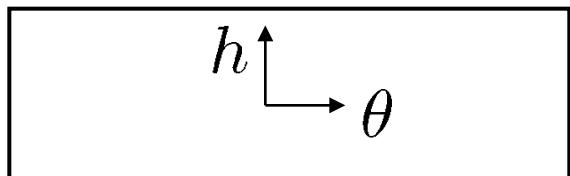
# Kogeto Dot 360 Camera for iPhone

---

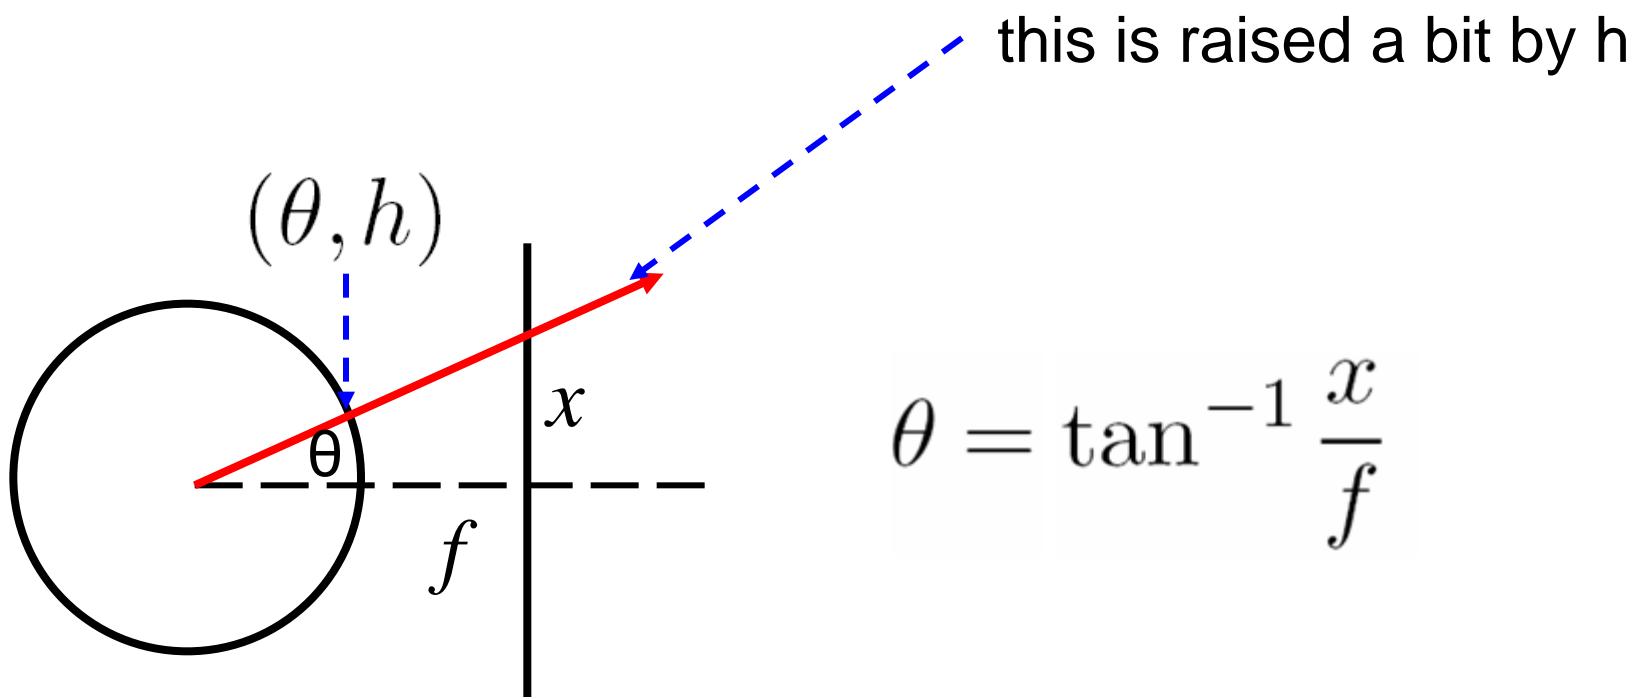
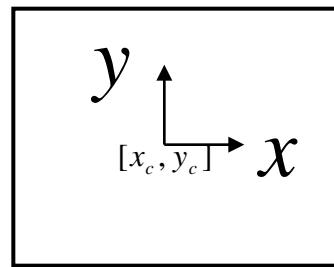


# Cylindrical projection

---

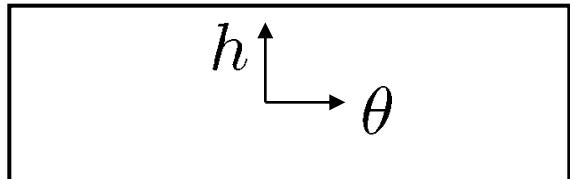


unwrapped cylinder

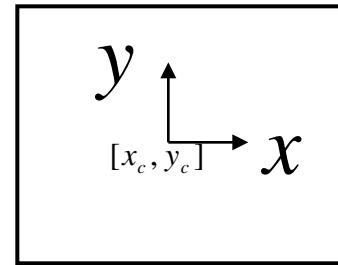


# Cylindrical projection

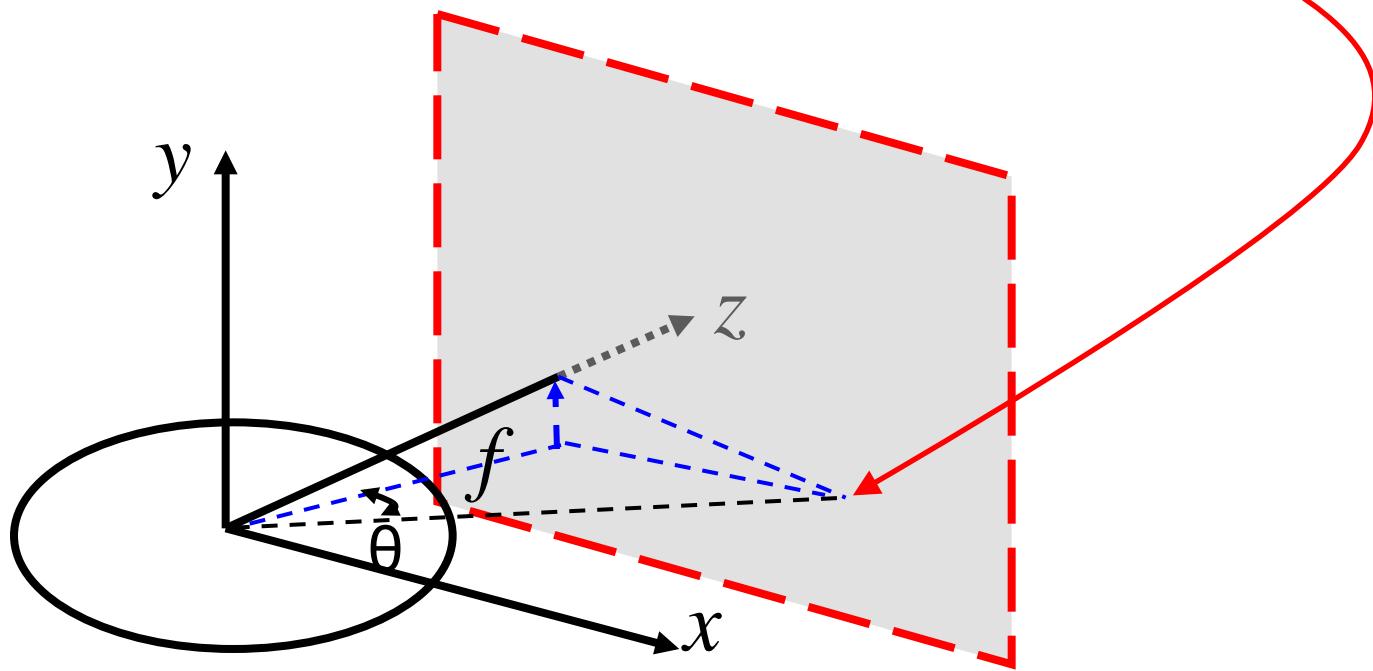
---



unwrapped cylinder

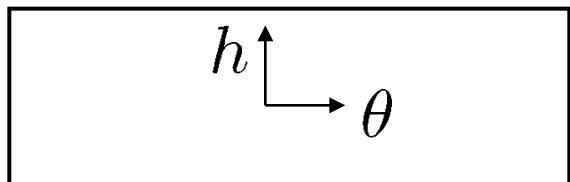


$\propto (x, y, f)$

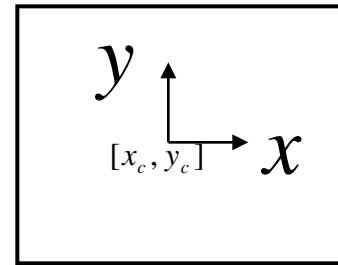


# Cylindrical projection

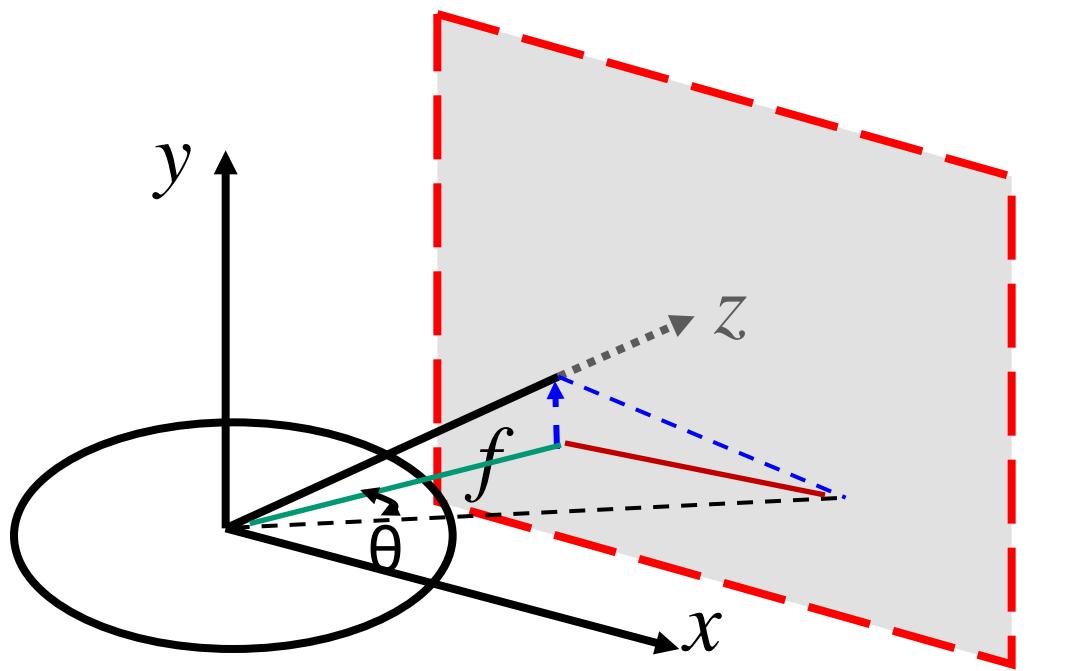
---



unwrapped cylinder

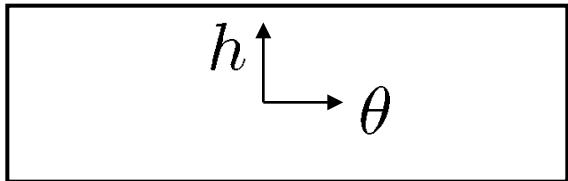


$$(\underline{\sin \theta}, h, \underline{\cos \theta}) \propto (x, y, f)$$

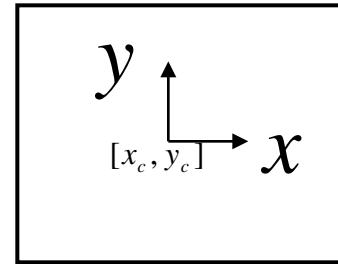


# Cylindrical projection

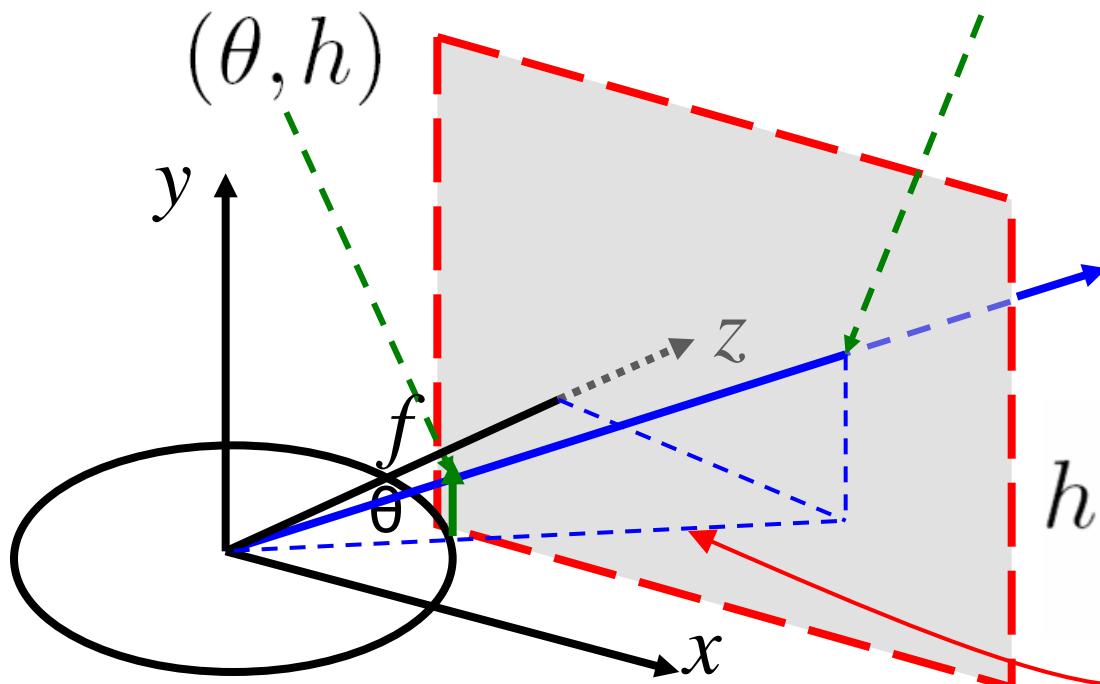
---



unwrapped cylinder

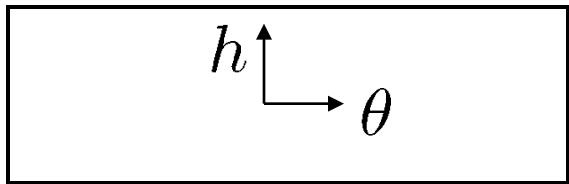


$$(\sin \theta, h, \cos \theta) \propto (x, y, f)$$

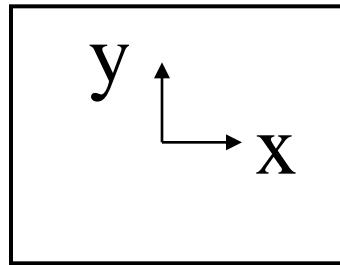


$$h = \frac{y}{\sqrt{x^2 + f^2}}$$

# Cylindrical Projection



unwrapped cylinder

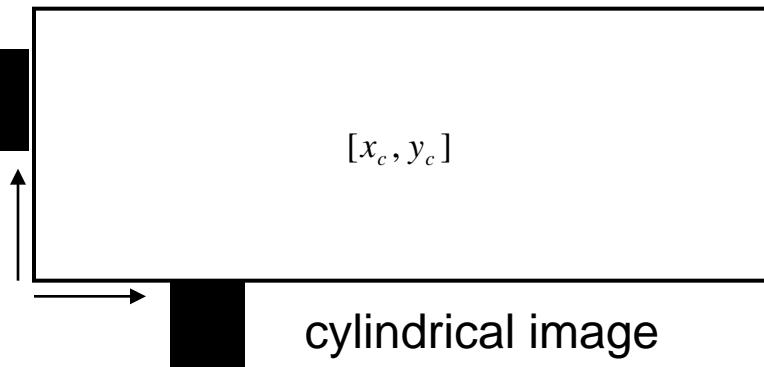
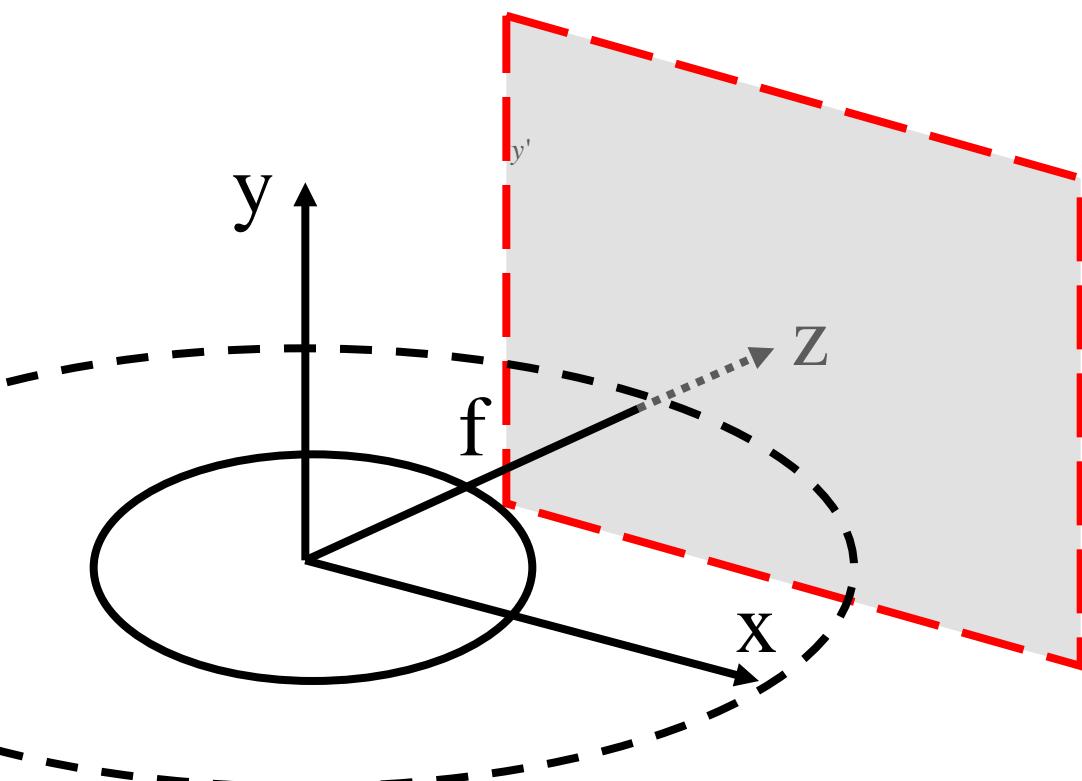


$[x_c, y_c]$

$$\tilde{x} = s\theta + x_c = s \tan^{-1} \frac{x}{f} + x_c$$

$$\tilde{y} = sh + y_c = s \frac{y}{\sqrt{x^2 + f^2}} + y_c$$

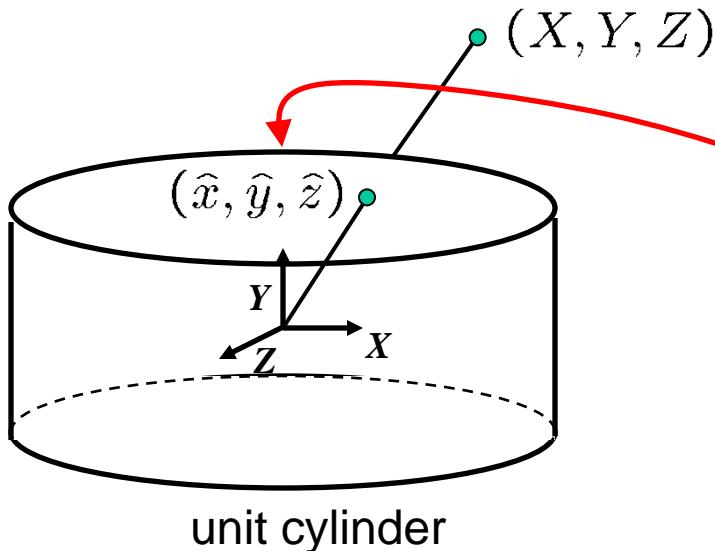
$s$  defines size of the final image,  
often convenient to set  $s = f$



cylindrical image

# Pipeline

---



Map 3D point/line  $(X, Y, Z)$  or  $Z=f$  onto cylinder

$$(\hat{x}, \hat{y}, \hat{z}) = \frac{1}{\sqrt{X^2+Z^2}}(X, Y, Z)$$

Convert to cylindrical coordinates

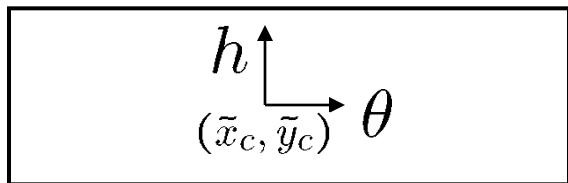
$$(\sin\theta, h, \cos\theta) = (\hat{x}, \hat{y}, \hat{z})$$

To get an “image”

Convert to cylindrical image coordinates

$$(\tilde{x}, \tilde{y}) = (f\theta, fh) + (\tilde{x}_c, \tilde{y}_c)$$

$s$  (or  $f$ ) defines size of the final image often convenient to set  $s = \text{camera focal length}$



unwrapped cylinder



cylindrical image

# Cylindrical image stitching

---



- What if you don't know the camera rotation?
  - Solve for the camera rotations
    - Note that a rotation of the camera is a **translation** of the cylinder

# Cylindrical image stitching

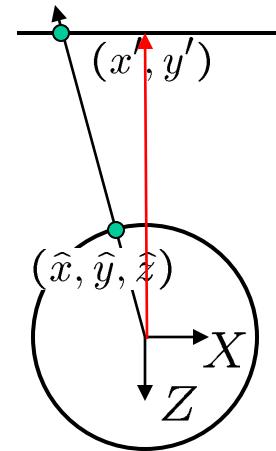
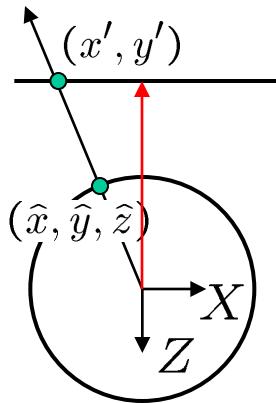
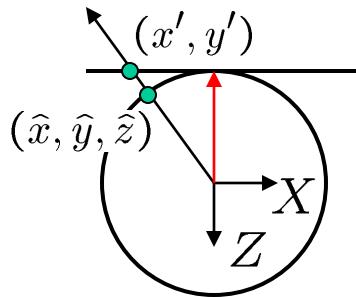
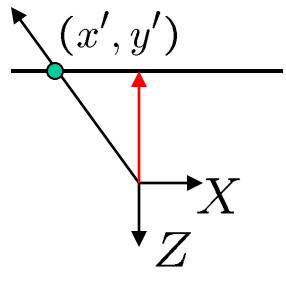
---



- What if you don't know the camera rotation?
  - Solve for the camera rotations
    - Remember we are not doing any perspective transformations!! Only translations

# Cylindrical reprojection

---



top-down view



Image 384x300



$f = 180$  (pixels)



$f = 280$

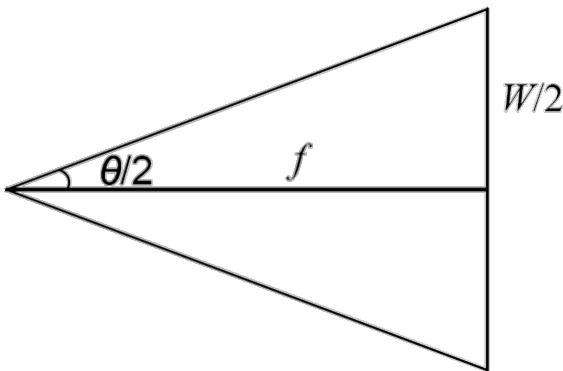


$f = 380$

# What's your focal length?

---

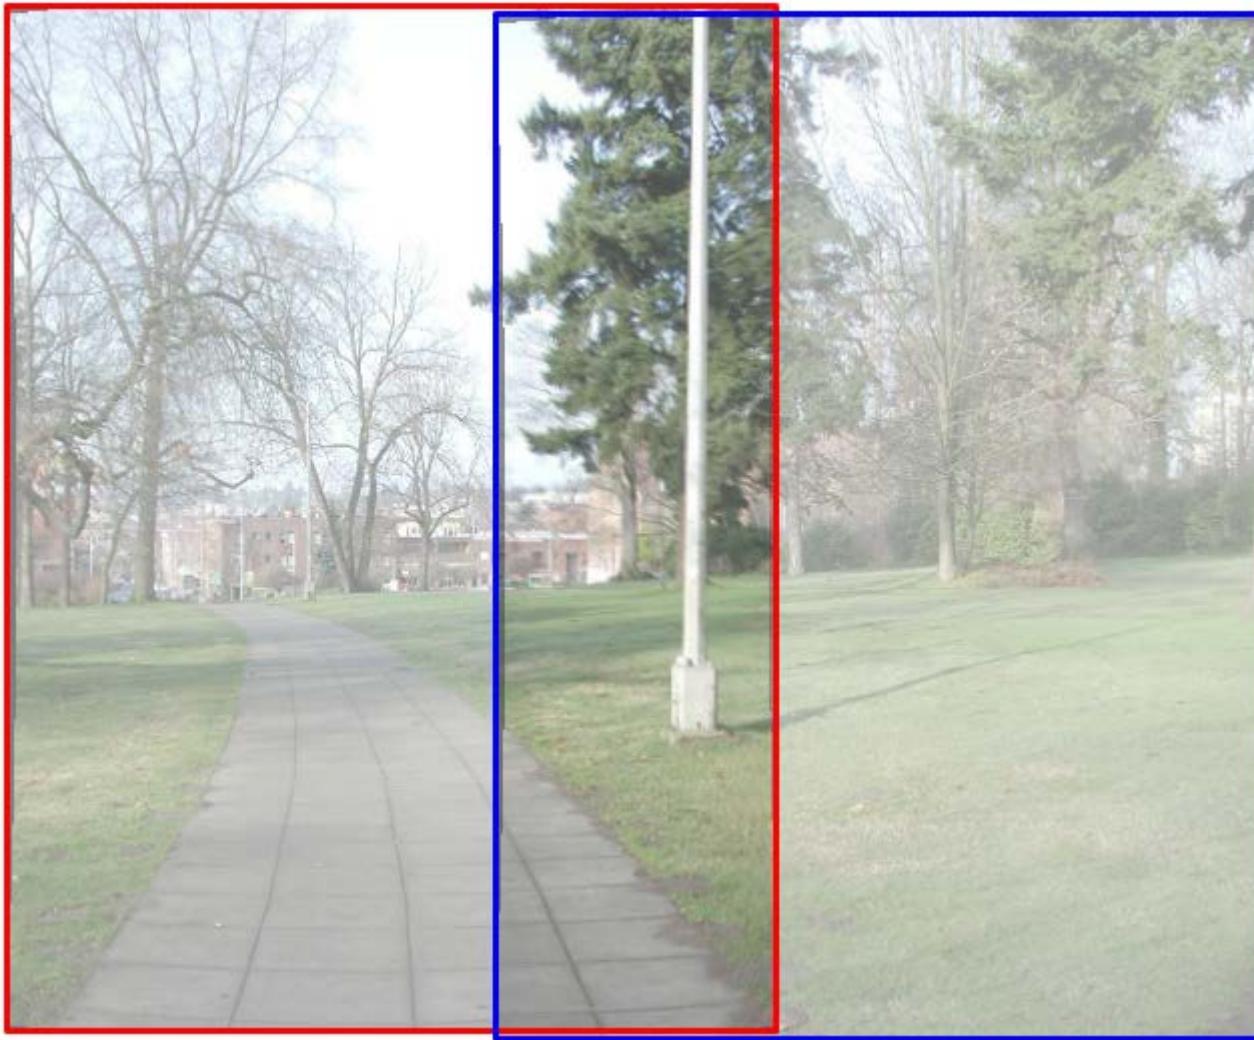
- Focal length is (highly!) camera dependent
  - Can get a rough estimate by measuring FOV:



- Can use the EXIF data tag (might not give the right info)
- Can use several images together and try to find  $f$  that would make them match
- Can use a known 3D object and its projection to solve for  $f$

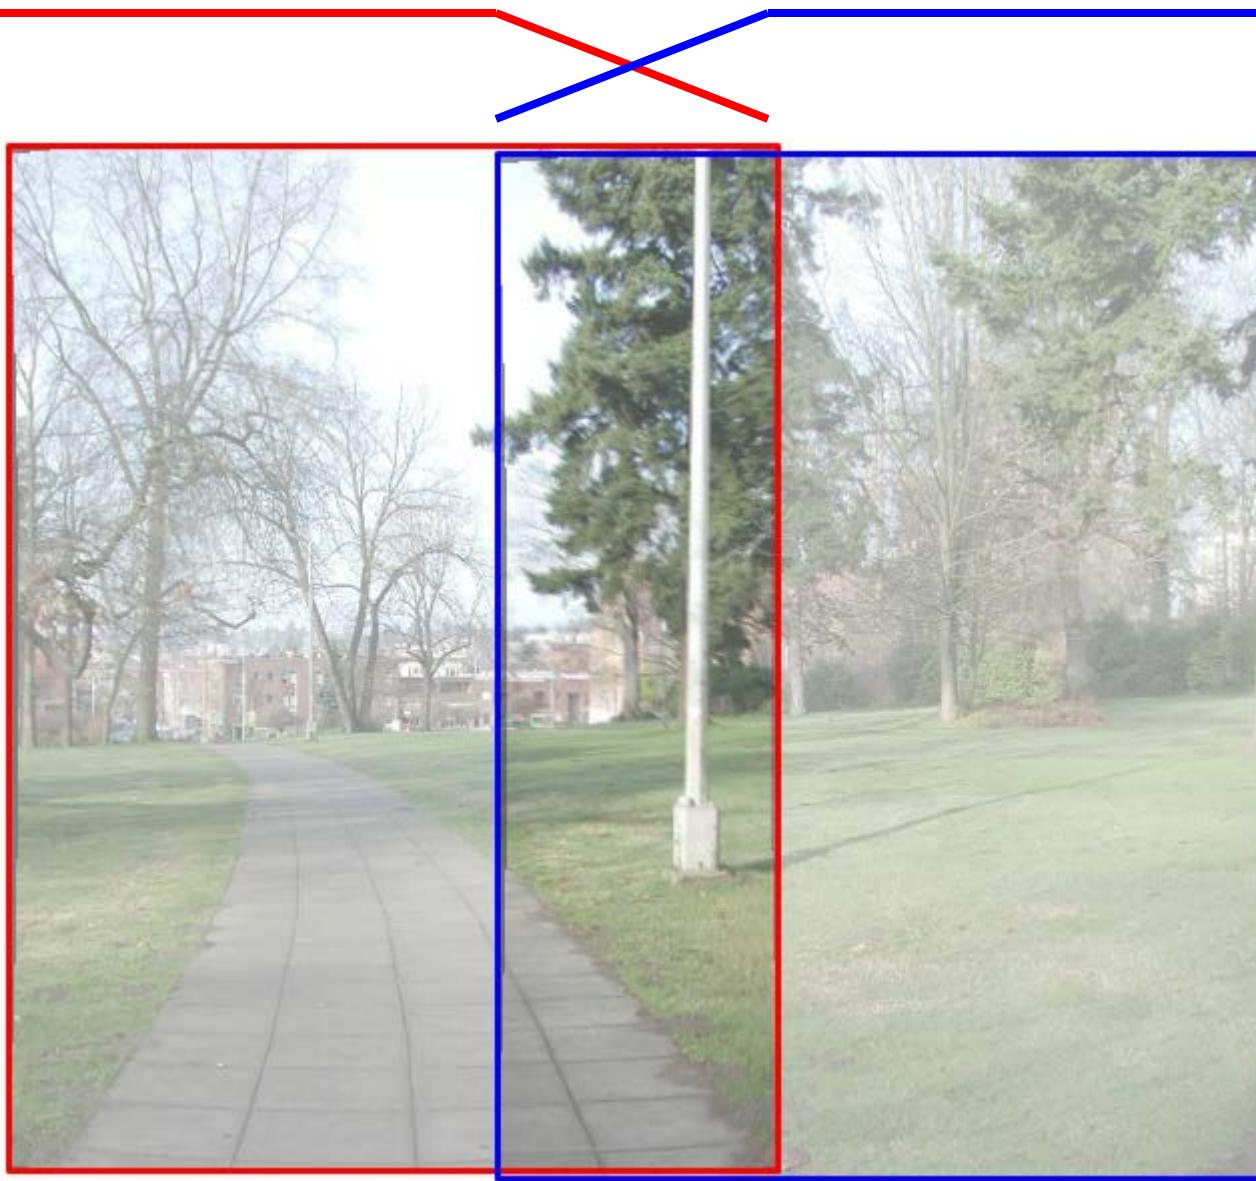
# Blending the seams

---



# Blending the seams

---



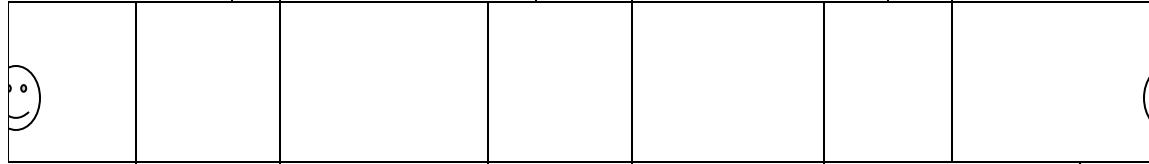
# Blending the seams

---



# Assembling the panorama

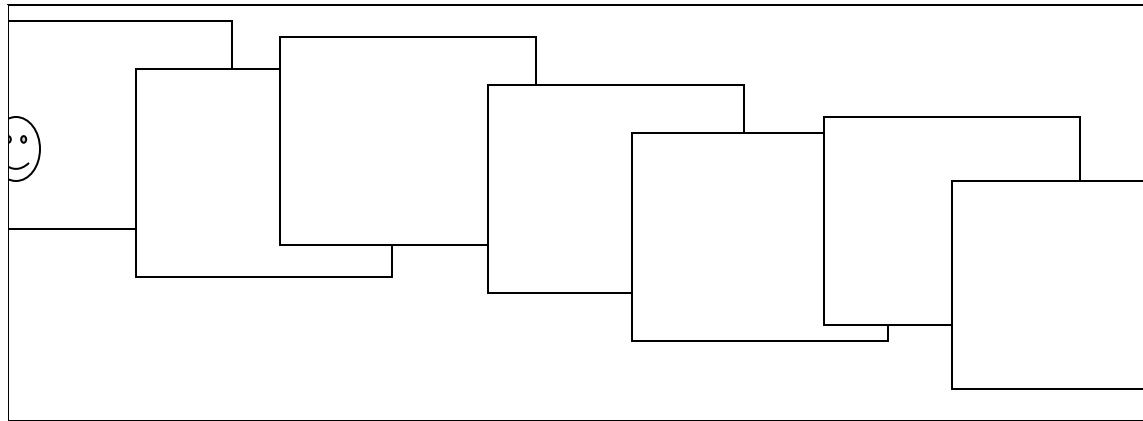
---



- Stitch pairs together, blend, then **crop**

# Problem: Drift

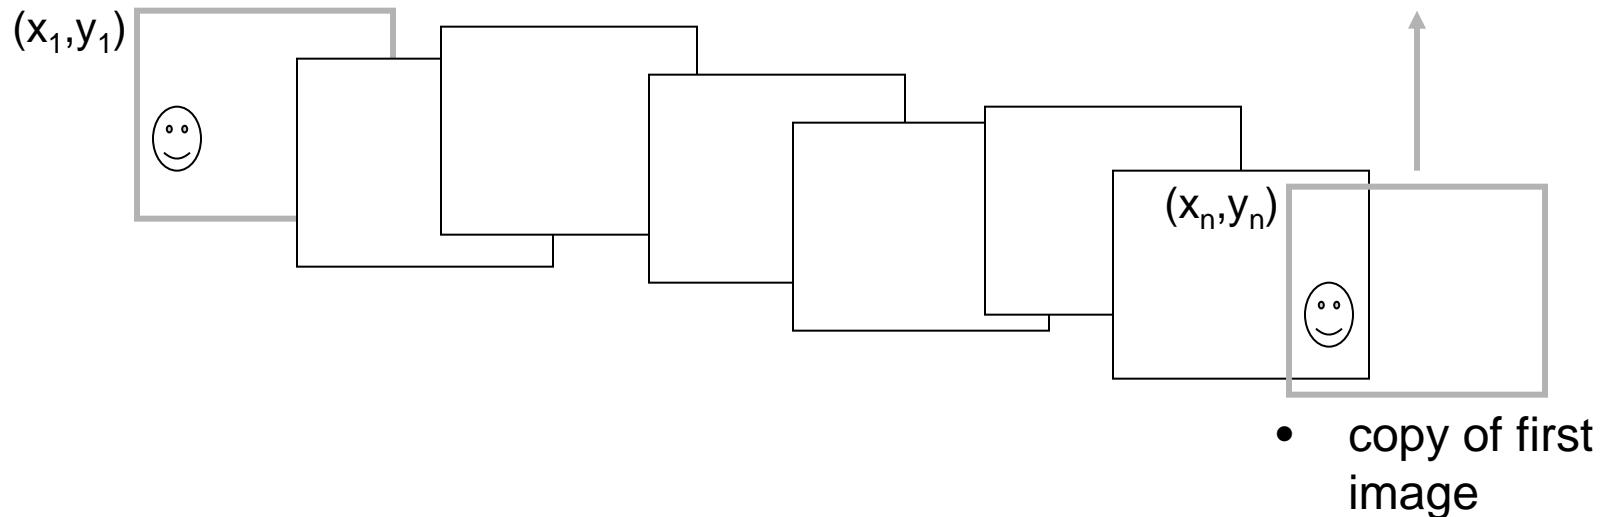
---



- Error accumulation
  - small errors accumulate over time

# Problem: Drift

---



- Error accumulation
  - small (vertical) errors accumulate over time
  - apply correction so that sum = 0

# End-to-end alignment and crop

---



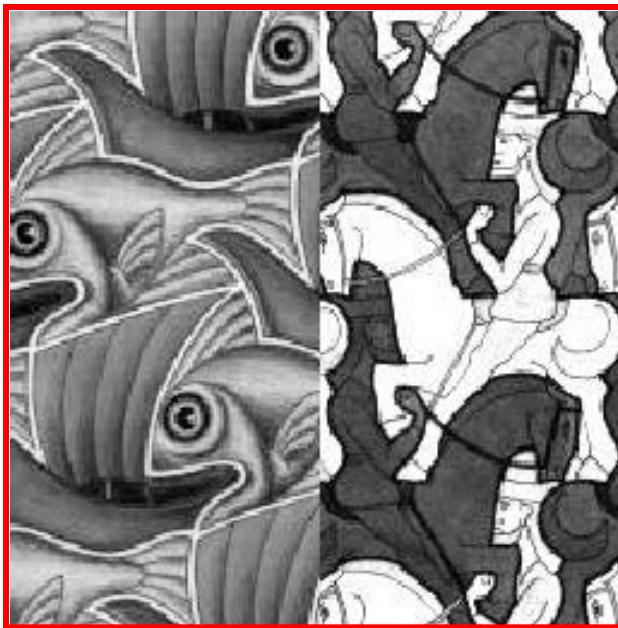
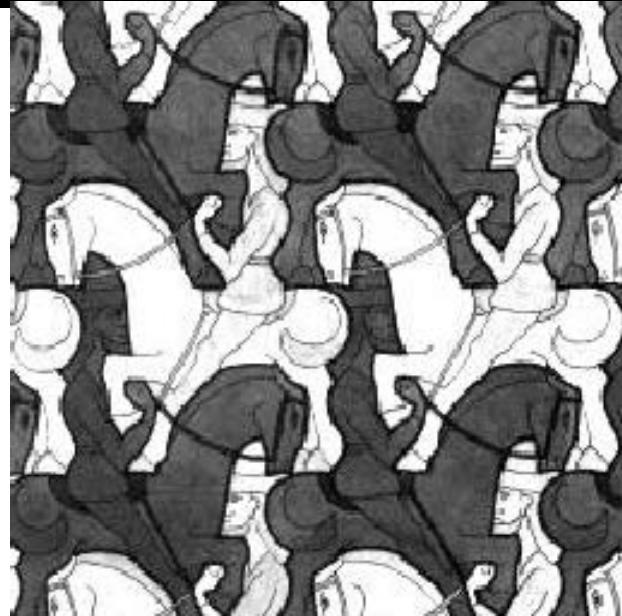
# Cylindrical panorama

---

1. Take pictures on a tripod (or handheld)
2. Warp to cylindrical coordinate
3. Compute pairwise alignments
4. Fix up the end-to-end alignment if needed
5. Blending
6. Crop the result

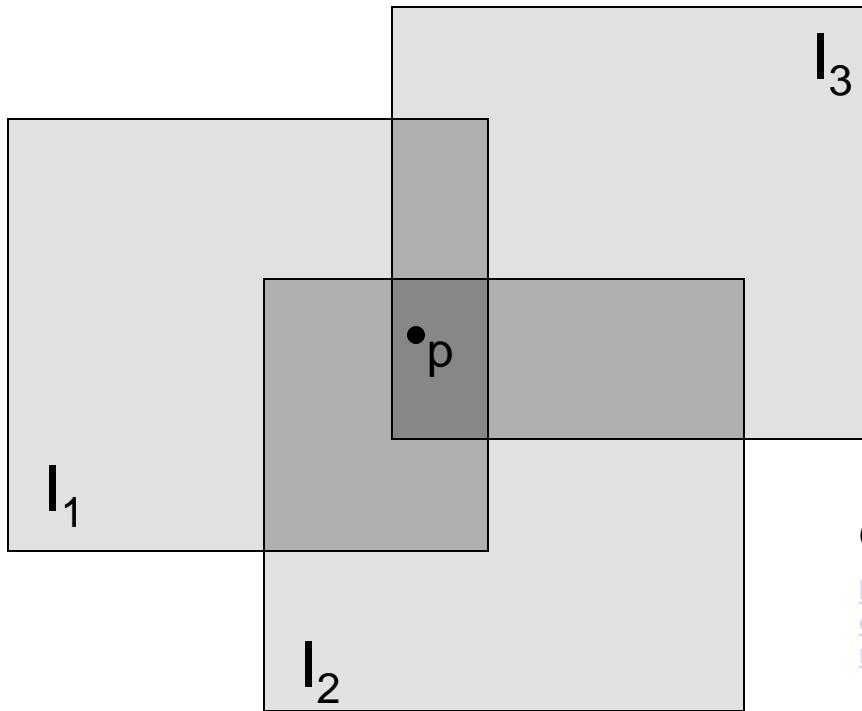
# Image Blending

---



# Alpha Blending

---



Optional: see Blinn (CGA, 1994) for details:

<http://ieeexplore.ieee.org/iel1/38/7531/00310740.pdf?isNumber=7531&prod=JNL&arnumber=310740&arSt=83&ared=87&arAuthor=Blinn%2C+J.F.>

Encoding blend weights

$$\text{Color at } p = \frac{(\alpha_1 R_1, \alpha_1 G_1, \alpha_1 B_1) + (\alpha_2 R_2, \alpha_2 G_2, \alpha_2 B_2) + (\alpha_3 R_3, \alpha_3 G_3, \alpha_3 B_3)}{\alpha_1 + \alpha_2 + \alpha_3}$$

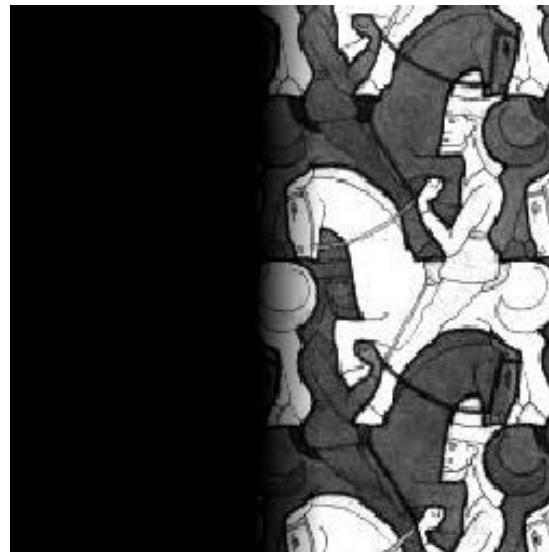
Implement this in two steps:

1. accumulate: add up the ( $\alpha$  premultiplied) RGB $\alpha$  values at each pixel
2. normalize: divide each pixel's accumulated RGB by its  $\alpha$  value

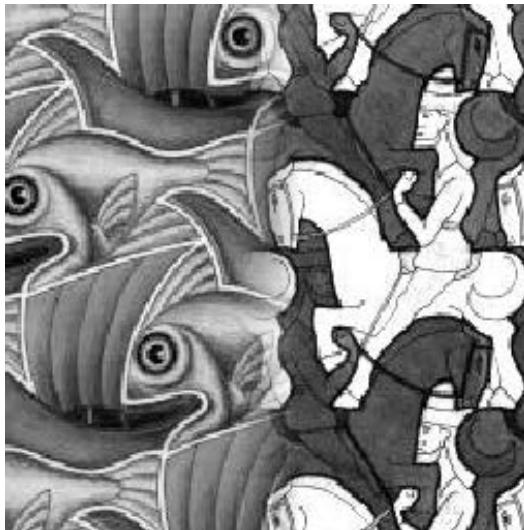
# Alpha Blending / Feathering



+



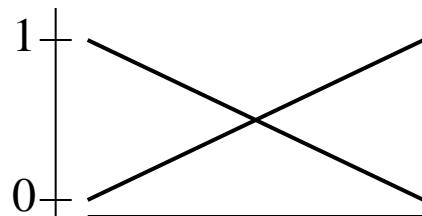
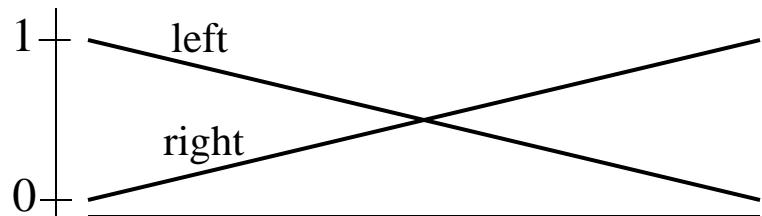
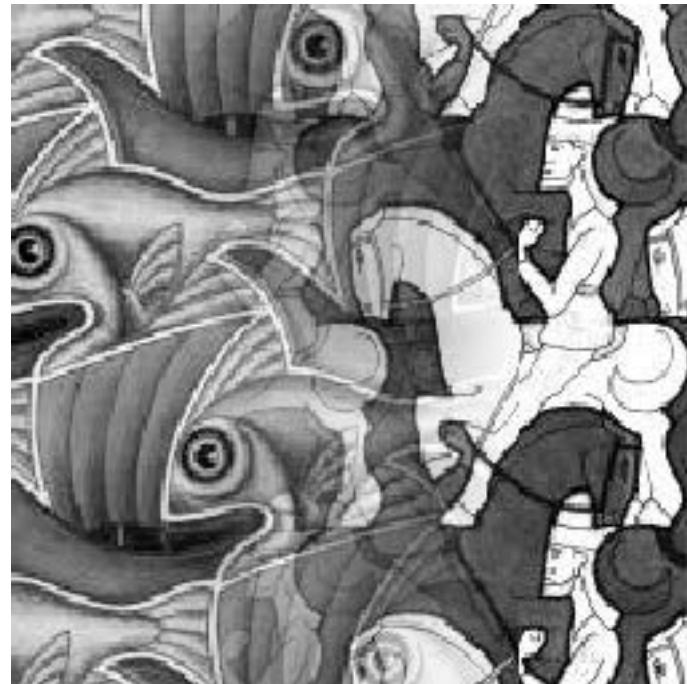
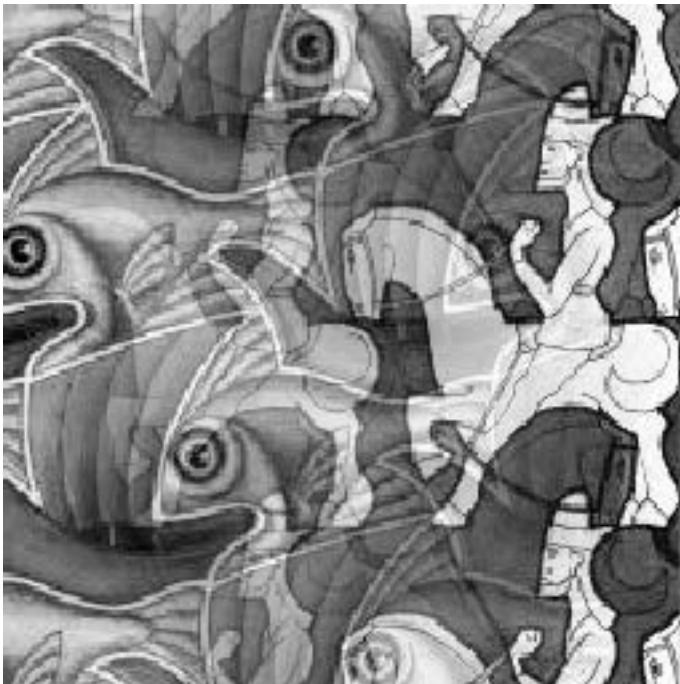
—



$$I_{\text{blend}} = \alpha I_{\text{left}} + (1-\alpha) I_{\text{right}}$$

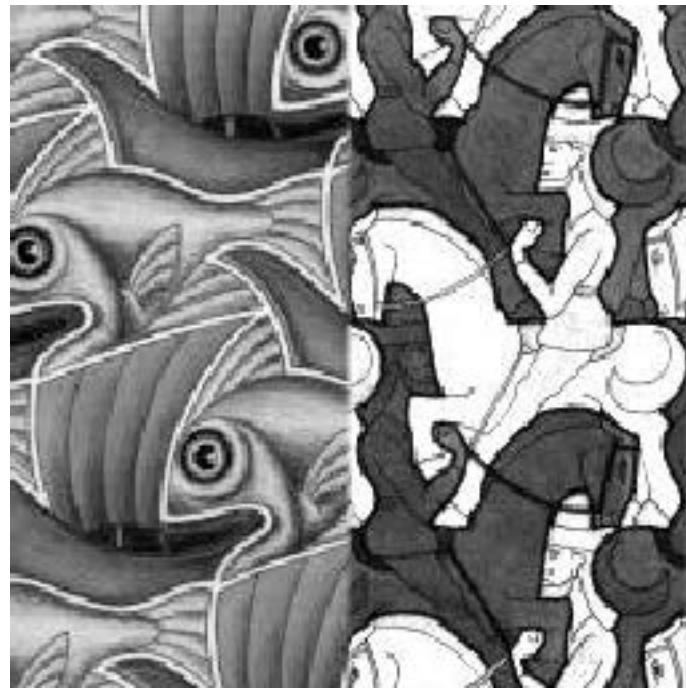
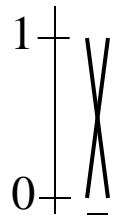
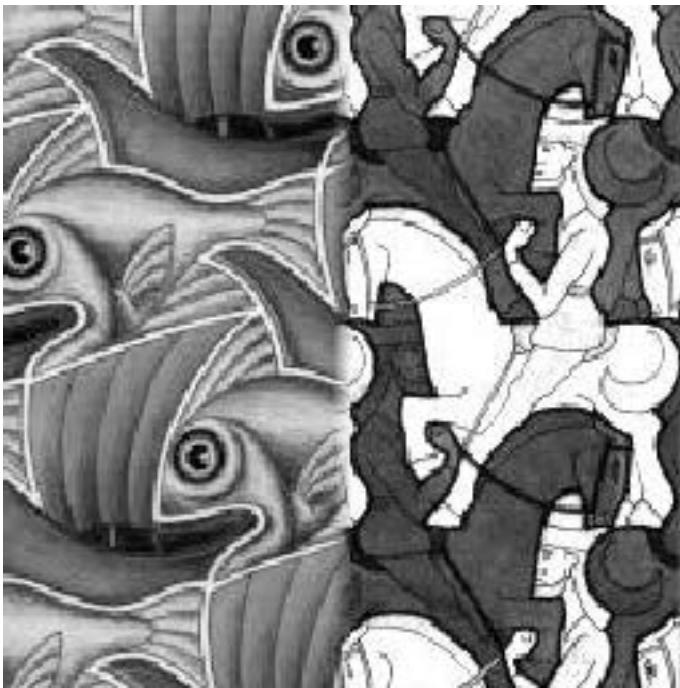
# Effect of window size

---



# Effect of window size

---

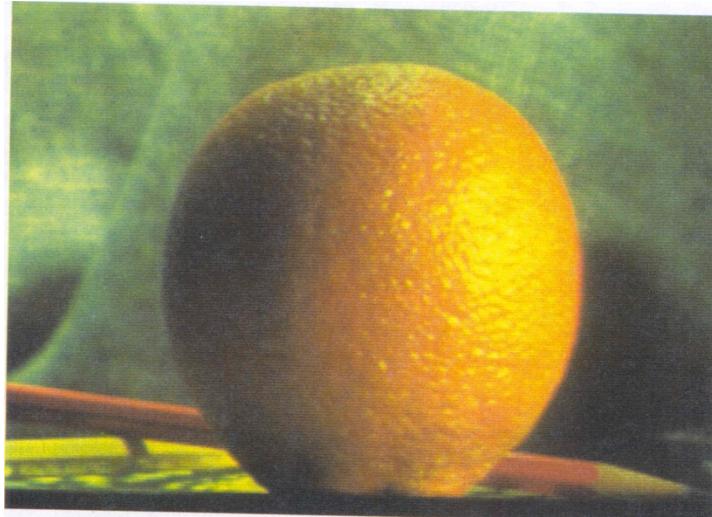


# Pyramid blending

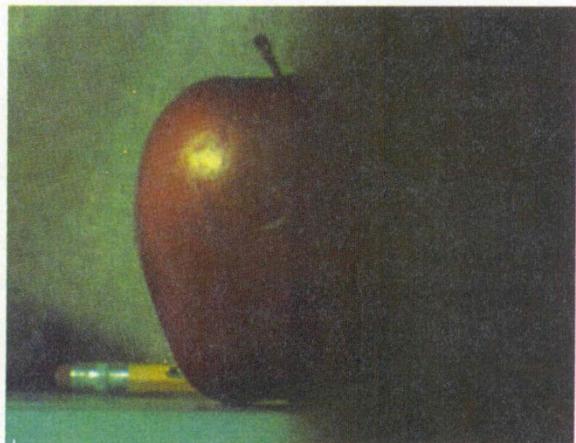
---



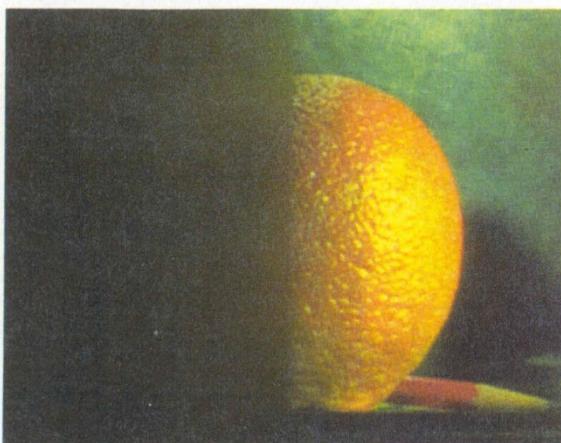
(a)



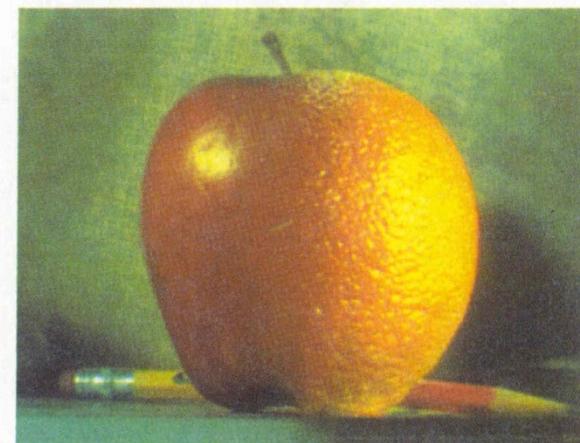
(b)



(d)



(h)



(l)

- Create a Laplacian pyramid, blend each level

- Burt, P. J. and Adelson, E. H., [A multiresolution spline with applications to image mosaics](#), ACM Transactions on Graphics, 42(4), October 1983, 217-236.

# Don't blend, CUT when there is motion!

---



Moving objects become ghosts

- So far we only tried to blend between two images. What about finding an optimal seam?

# Davis, 1998

---

- Segment the mosaic
  - Single source image per segment!!!
  - Avoids artifacts along boundaries
    - Dijkstra's algorithm



# Davis, 1998

---

- Segment the mosaic
  - Single source image per segment
  - Avoid artifacts along boundaries
    - Dijkstra's algorithm



# Poisson Image Editing (not panoramas)



sources/destinations



cloning



seamless cloning

- For more info: Perez et al, SIGGRAPH 2003

- [http://research.microsoft.com/vision/cambridge/papers/perez\\_siggraph03.pdf](http://research.microsoft.com/vision/cambridge/papers/perez_siggraph03.pdf)

# Some panorama examples

---



Microsoft Lobby: <http://www.acm.org/pubs/citations/proceedings/graph/258734/p251-szeliski>

# Some panorama examples

---

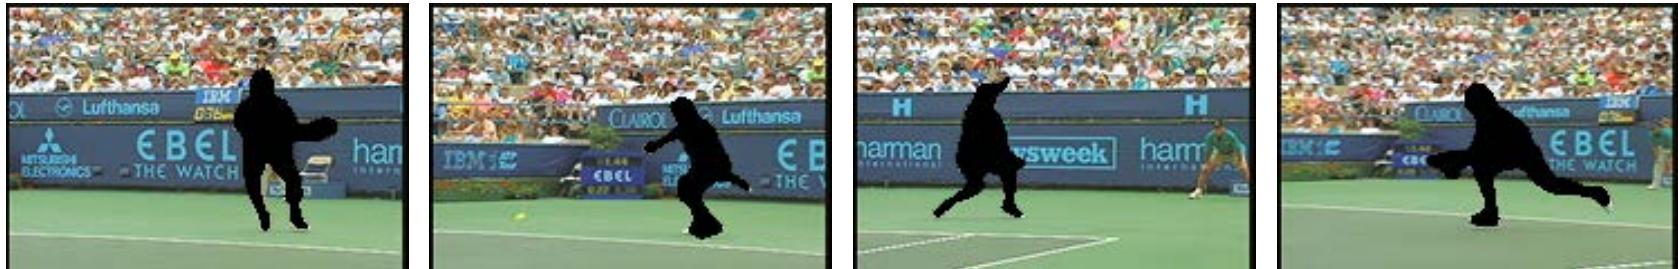


What's inside your refrig?

<http://www.cs.washington.edu/education/courses/cse590ss/01wi/>

# Video compression?

---



# Magic: ghost removal

---

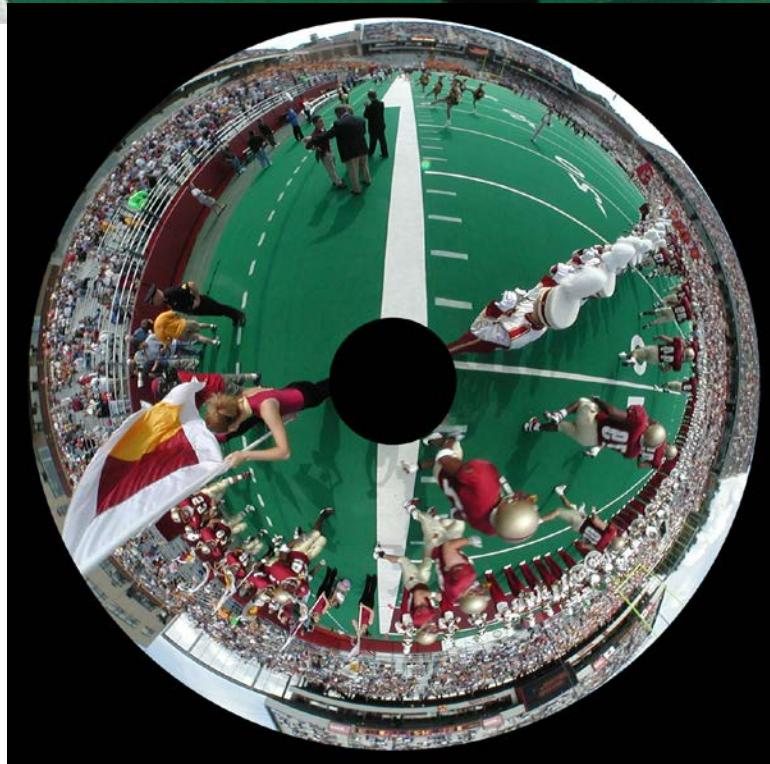


# Magic: ghost removal

---



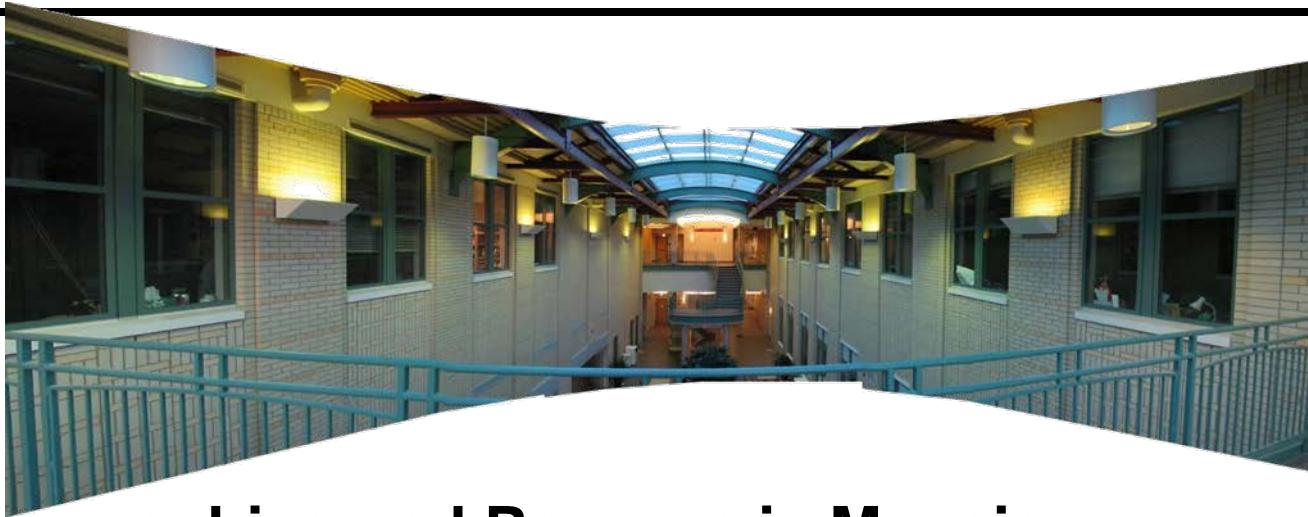
M. Uyttendaele, A. Eden, and R. Szeliski.  
*Eliminating ghosting and exposure artifacts in image mosaics.*  
CVPR 2001





# Programming Project #2

---



- **Homographies and Panoramic Mosaics**
- Capture photographs (and possibly video)
  - can check out cameras and/or tripods
- Compute homographies (define correspondences)
  - will need to figure out how to setup system of eqs.
- (un)warp an image (undo perspective distortion)
- Produce a few panoramic mosaics (with blending)
- Do some of the Bells and Whistles

# When does a Video lead to a Good Panorama?

---

- Good homographies between frames
- Individual images have good quality
- Result has a wide field of view

Competing interests:

- More frames give wider field of view
- More frames give more accumulated error

# In defense of 8 point algorithm (Hartley)

---

Observation: Linear estimation of projective transformation parameters from point correspondences often suffer from poor “conditioning” of the matrices involved. This means the solution is sensitive to noise in the points (even if there are no outliers).

Solution: To get better answers, precondition the matrices by performing a normalization of each point set by:

- translating center of mass to the origin
- scaling so that average distance of points from origin is  $\sqrt{2}$ .